

[Goto table of contents](#)



Exploring WebAssembly for versatile plugin systems through the example of a text editor

im Studiengang TINF22IT1

an der *Duale Hochschule Baden-Württemberg Mannheim*

von

Name, Vorname: Hartung, Florian

Abgabedatum: 15.04.2025

Bearbeitungszeitraum: 15.10.2024 - 15.04.2024

Matrikelnummer, Kurs: 6622800, TINF22IT1

Wiss. Betreuer*in der Dualen Hochschule: Gerhards, Holger, Prof. Dr.

Erklärung zur Eigenleistung

Ich versichere hiermit, dass ich meine Projektarbeit mit dem

THEMA

Exploring WebAssembly for versatile plugin systems through the example of a text editor

selbstständig verfasst und keine anderen als die angegebenen Quellen und Hilfsmittel benutzt habe. Ich versichere zudem, dass die eingereichte elektronische Fassung mit der gedruckten Fassung übereinstimmt.

Ort, Datum

Unterschrift

Abstract

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magnam aliquam quaerat voluptatem. Ut enim aequae doleamus animo, cum corpore dolemus, fieri tamen permagna accessio potest, si aliquod aeternum et infinitum impendere malum nobis opinemur. Quod idem licet transferre in voluptatem, ut postea variari voluptas distinguique possit, augeri amplificarique non possit. At etiam Athenis, ut e patre audiebam facete et urbane Stoicos irridente, statua est in quo a nobis philosophia defensa et collaudata est, cum id, quod maxime placeat, facere possimus, omnis voluptas assumenda est, omnis dolor repellendus. Temporibus autem quibusdam et aut officiis debitis aut rerum necessitatibus saepe eveniet, ut et voluptates repudiandae sint et molestiae non recusandae. Itaque earum rerum defuturum, quas natura non depravata desiderat. Et quem ad me accedis, saluto: 'chaere,' inquam, 'Tite!' lictores, turma omnis chorusque: 'chaere, Tite!' hinc hostis mi Albucius, hinc inimicus. Sed iure Mucius. Ego autem mirari satis non queo unde hoc sit tam insolens domesticarum rerum fastidium. Non est omnino hic docendi locus; sed ita prorsus existimo, neque eum Torquatum, qui hoc primum cognomen invenerit, aut torquem illum hosti detraxisse, ut aliquam ex eo est consecutus? – Laudem et caritatem, quae sunt vitae.

Zusammenfassung

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magnam aliquam quaerat voluptatem. Ut enim aequaleamur animo, cum corpore dolemus, fieri tamen permagna accessio potest, si aliquod aeternum et infinitum impendere malum nobis opinemur. Quod idem licet transferre in voluptatem, ut postea variari voluptas distinguique possit, augeri amplificarique non possit. At etiam Athenis, ut e patre audiebam facete et urbane Stoicos irridente, statua est in quo a nobis philosophia defensa et collaudata est, cum id, quod maxime placeat, facere possimus, omnis voluptas assumenda est, omnis dolor repellendus. Temporibus autem quibusdam et aut officiis debitis aut rerum necessitatibus saepe eveniet, ut et voluptates repudiandae sint et molestiae non recusandae. Itaque earum rerum defuturum, quas natura non depravata desiderat. Et quem ad me accedis, saluto: 'chaere,' inquam, 'Tite!' lictores, turma omnis chorusque: 'chaere, Tite!' hinc hostis mi Albucius, hinc inimicus. Sed iure Mucius. Ego autem mirari satis non queo unde hoc sit tam insolens domesticarum rerum fastidium. Non est omnino hic docendi locus; sed ita prorsus existimo, neque eum Torquatum, qui hoc primus cognomen invenerit, aut torquem illum hosti detraxisse, ut aliquam ex eo est consecutus? – Laudem et caritatem, quae sunt vitae.

Contents

1 Introduction	1
1.1 Motivation & problem statement	1
1.2 Research question	1
1.3 Methodology	1
2 Fundamentals	2
2.1 Instruction set architectures	2
2.2 WebAssembly	2
2.2.1 Overview	2
2.2.2 Execution model and lifecycle	5
2.2.3 Design goals	6
2.2.4 Challenges & Limitations	9
2.2.5 WebAssembly System Interface (WASI)	9
2.2.6 WebAssembly Component Model	9
2.3 Plugin systems	9
3 Criteria for good plugin systems	10
3.1 Definition of criteria	10
3.1.1 Performance	10
3.1.2 Plugin size	11
3.1.3 Plugin isolation	13
3.1.4 Plugin portability	14
3.1.5 Plugin language interoperability	16
3.2 A technology comparison between existing plugin systems	17
3.2.1 Choice of technologies & projects	17
3.2.2 Evaluations of technologies & projects	20
3.2.3 Summary	24
4 WebAssembly for plugin systems	25
4.1 Overview of basic plugin system architecture	25
4.2 Evaluation of requirements	25

4.3 Evaluation of interface-specific requirements	25
4.3.1 Without a standardized interface	25
4.3.2 WebAssembly Component Model	25
4.3.3 WebAssembly Component Model + WebAssembly System Interface (WASI)	26
4.3.4 Custom serialization format (JSON, XML, Protobuf)	26
4.4 Plugin systems already using WebAssembly today	26
4.5 Summarized evaluation for WebAssembly	27
5 Proof of concept: Implementing a WebAssembly plugin system for a text editor	28
5.1 Requirements	28
5.1.1 Functional requirements	28
5.1.2 Non-functional requirements	28
5.2 System Architecture	28
5.3 Implementation	28
5.4 Evaluation & Results	28
6 Results & Discussion (2 pages)	29
7 Outlook	30
Bibliography	31

List of figures

Figure 1: A multilevel computer system running WebAssembly code	3
Figure 2: Flowchart of the lifecycle of a WebAssembly module	5
Figure 3: An overview over all WebAssembly opcodes	7

List of tables

Table 1: Technology comparison matrix of existing technologies & projects	24
---	----

1 Introduction

1.1 Motivation & problem statement

1.2 Research question

Is WebAssembly the best technology choice for designing versatile plugin systems for text editors?

1.3 Methodology

2 Fundamentals

This section introduces theoretical and technical fundamentals used in this work. First it covers the definition of an instruction set architecture. Then it gives an overview over WebAssembly, its features, challenges, limitations and extensions. Finally, plugin systems are explained as a software architecture model.

2.1 Instruction set architectures

In his book “Structured Computer Organization” Tanenbaum defines an instruction set architecture (ISA) as a level in a multilayered computer system[1, sec. 1.1.2]. The ISA level defines a machine language with a fixed set of instructions. According to Tanenbaum the ISA level then acts as a common interface between the hardware and software. This allows software in the form of ISA instructions to manipulate the hardware. Software written in a higher level machine language (Assembly, C, Java, ...) can not be executed directly by the hardware. Instead higher level machine codes are compiled to ISA machine code or interpreted by a program, that is present in ISA machine code itself [1, sec. 1.1.2].

TODO: move Figure 1 here and explain basics of multi-level systems here separately

2.2 WebAssembly

2.2.1 Overview

WebAssembly (Wasm) is a stack-based ISA for a portable, efficient and safe code format. Originally it was designed by engineers from the four major vendors to enable high-performance code execution on the web [2]. However it is also becoming increasingly interesting for researchers and developers in non-web contexts. Some examples are avionics for Wasm’s safe and deterministic execution [3], distributed computing for its portability and migratability [4] or embedded systems for its portability and safety [5].

What is special about Wasm is that it is a *virtual* ISA [6, sec. 1.1.2]. There is no agreed-upon definition for a virtual ISA, however the term *virtual* can be assumed to refer to an ISA that is running in a virtualized environment on a higher level in

2.2 WebAssembly




Figure 1: *TODO: This figure still needs some improvements*. A multilevel computer system running Wasm code. Based on Figure 1-2 in [1]

a multilevel computer¹. We call this virtualized environment the **host environment** (used by the specification[6, sec. 1.2.1]) or the **WebAssembly runtime** (used by most technical documentation).

If one considers a system running Wasm code as a multi-level computer system, the Wasm runtime can be modeled as a separate layer. Figure 1 shows a multi-level computer system based on Tanenbaum’s model [1, sec. 1.1.2]. Here each level is executed by logic implemented in the next lower level either through compilation or interpretation. The digital logic level itself only exists in the form of individual gates, consisting of transistors and tracks on the processors’ chip. This level runs the next microarchitecture and ISA levels, which are also often implemented directly in hardware. The ISA level then provides a fixed set of instructions for higher levels to use. Operating systems build on top of this and provide another level for user space programs, which exist on the assembly language level. Then there are problem-oriented languages such as C, C++ or Rust, which are specifically made for humans to write code in[1, sec. 1.1.2].

One program written in a problem-oriented language is the Wasm runtime, which itself is a layer here. Its task it to interpret or compile higher-level Wasm code to lower-level problem-oriented or even the assembly language level. However for this work all layers starting with the Wasm runtime level until the digital logic level can be seen as a single hardware specific layer called the *Wasm runtime environment*.

¹Projects that try to execute Wasm directly exist. One example is the discontinued wasmachine project, which tried executing Wasm on FPGAs: <https://github.com/piranna/wasmachine>

2.2 WebAssembly

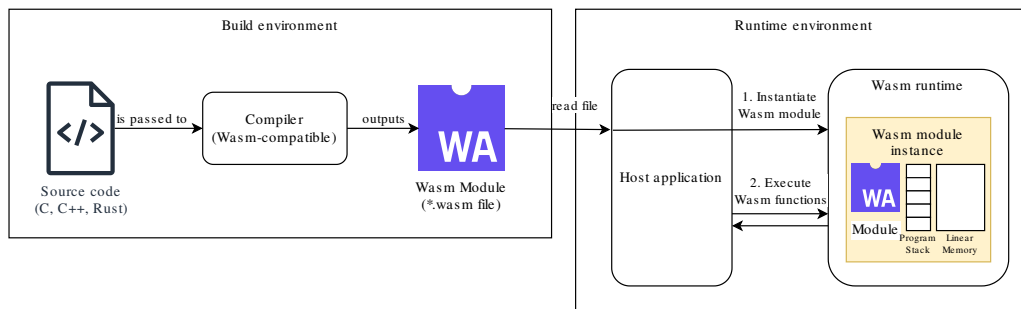


Figure 2: Flowchart for the creation and execution of a Wasm module from a higher-level language *TODO: This figure is still too complicated*

WebAssembly code can be written by hand, just as one could write traditional ISA instructions (think of writing x86 machine code) by hand. Even though it is possible, it would not be efficient to write useful software systems in Wasm because the language is too simple. For example it provides only a handful of types: Signed/unsigned integers, floating point numbers, a 128-bit vectorized number and references to functions/ host objects.

WebAssembly also does not provide any way to specify memory layouts as it can be done in higher-level languages with structs, classes, records, etc. Instead it provides most basic features and instructions, which exist on almost all modern computer architectures like integer and floating point arithmetic, memory operations or simple control flow constructs. This is by design, as WebAssembly is more of a compilation target for higher level languages². Those higher level languages can then build upon Wasm's basic types and instructions and implement their own abstractions like memory layouts or control flow constructs on top. This is analogous to the non-virtual ISA machine code in a conventional computer, which also acts as a compilation target for most low-level languages such as Assembly, C, or Rust. Nowadays there are compilers for most popular languages already. Some of which are `clang` for C/C++, `rustc` for Rust or the official Zig and Haskell compilers. However most compilers are still being actively worked on and improved over time to support the latest Wasm proposals.

²<https://webassembly.org/>

2.2 WebAssembly

2.2.2 Execution model and lifecycle

Figure 2 shows the different stages a Wasm module goes through in its lifecycle. Its lifecycle starts with a developer writing source code. This source code can be written in an arbitrary programming language such as C, C++ or Rust. These languages are often used because they have compilers that support Wasm. The compiler then compiles this source code to a *Wasm module*. This step requires a compiler that support Wasm as a compilation target. Many modern compilers such as `clang` or `rustc` use the LLVM³ project for optimization and code generation. These compilers compile source code to LLVM’s intermediate representation (LLVM IR) and pass this to LLVM. LLVM can then perform optimizations and compile this IR to any compilation target, which can be selected by choosing a LLVM backend. One such LLVM backend targets Wasm. This way the Wasm-specific compiler logic can be implemented once within LLVM, and any compiler using LLVM can then target Wasm with minimal additional effort.

The compiled Wasm module exists in the form of a `.wasm` file. It is fully self-contained and unlike binaries it cannot depend on any dynamic libraries at runtime. It consists of different ordered sections, each with their own purpose: Some example sections contain function signatures, data segments, import- and export definition or actual Wasm instructions for each function.

The previously generated Wasm module can then be transferred to any target device or platform providing a Wasm runtime. That specific Wasm runtime is hardware-specific and not portable, unlike the Wasm module. The Wasm runtime is able to parse the received Wasm module file, instantiate a *Wasm instance* from it and provide an Application Programming Interface (API) for interaction with the Wasm instance. APIs can differ from one Wasm runtime to another. Some runtimes exist as standalone programs that can run Wasm modules comparable to how native binaries can be executed. Others are in the form of libraries, that can only be used from a host application to embed a Wasm runtime into them. These Wasm runtime libraries often provide common operations to the host application like calling Wasm functions, reading and writing operations for Wasm memories, linking mechanisms between Wasm modules, exposing host-defined functions for Wasm instances to call, etc.

³According to its official website the acronym *LLVM* was once short for *Low Level Virtual Machine*, however now it has “grown to be an umbrella project” referred to simply as the LLVM project. <https://llvm.org/>

2.2 WebAssembly

In a web context a server might provide this Wasm module to the client's browser, which usually comes with a Wasm runtime. Detailed information on the current status of Wasm support for internet browsers can be viewed at <https://caniuse.com/wasm>. A concrete example is Ebay using Wasm for their barcode scanner algorithm to achieve higher performance⁴.

For distributed computing a compiled Wasm module could be distributed among multiple different nodes regardless of their platform and architecture... *TODO*

2.2.3 Design goals

Wasm was designed with certain design goals in mind. This section presents the design goals relevant for this work according to the official Wasm specification [6, sec. 1.1.1]. Each design goal is accompanied by related information from various papers and articles for a deeper understanding of each goal.

2.2.3.A Fast

Wasm is designed to be fast both during startup and execution [6, sec. 1.1.1]. Startup time is mostly optimized through the structure of the Wasm bytecode format, which is optimized for fast parsing and compiling of Wasm code. A Wasm module in its binary format consists of 11 different sections⁵ that may only show up in a fixed order. Another example is the absence of backwards jump (except for the `loop` statement), which allows the use of faster one-pass compilers.

During runtime Wasm can also achieve near-native performances... *TODO*

An additional property of Wasm bytecode is that it can be either compiled, interpreted or just-in-time compiled by a runtime. This allows users to customize Wasm runtimes for their specific needs and use cases. For example one could choose compilation to achieve faster execution at the cost of slow startup times due to compilation. On the other end users might prefer interpretation where execution speeds are less of a priority and fast startup times or relocatable runtime instances (see Section 2.2.3.G) are needed.

⁴<https://innovation.ebayinc.com/tech/engineering/webassembly-at-ebay-a-real-world-use-case/>

⁵This section count excludes custom sections, which are optional and may only add additional information on top, such as debug information.

2.2 WebAssembly

TODO: opcode table

Figure 3: All WebAssembly opcodes. Opcodes for proposals are encoded by specific marker bytes, which indicate that the following opcodes are to be interpreted as different instructions.

2.2.3.B Safe

TODO:

- *safe by default*
- *sandboxed*
 - *Memory access is safe: Wasm is stack based & provides a linear memory, which can only be accessed through indices with bounds checks on every access.*
 - *Can only interact with host environment through functions that are explicitly exposed by the host*
- *Because of its simplicity: easier to implement a minimal working runtime than it is for higher level languages like C, C++, Java, Python, ...*
 - *This again reduces risk of (safety-critical) bugs*
- *Properties like portability and safety are especially important in the context of the web, where untrusted software from a foreign host is executed on a client's device.*
- *This makes WASM, a sandboxed and fast execution environment, interesting for safety-critical fields like avionics (TODO: ref?, automotive (TODO: ref <https://oxidos.io/>), TODO: what else?)*
- *Note: Applications inside Wasm can still corrupt their own memory.*

2.2.3.C Portable

Wasm is designed to be able to be portable for a lot of different hardwares and platforms.

TODO:

- *minimal set of opcodes (172), that exist on all architectures (TODO: insert screenshot of opcode table (+ proposals for more additional instructions like SIMD, atomics, etc.))*

Independence of hardware:

- *Desktop architectures*
- *Mobile device architectures*
- *Embedded system architectures*

2.2 WebAssembly

Independence of platform:

- *Browsers*
- *Other environments which only require some kind of a Wasm runtime*

2.2.3.D Independence of language

TODO:

- *Not designed for a specific language, programming model or object model[6, sec. 1.1.1]*
- *Should act as a compilation target for all kinds of higher-level machine languages*
- *Some Examples for languages that can be used at the time of writing are: TODO*

2.2.3.E Compact

TODO:

- *Wasm bytecode representation format should be compact.*
- *Smaller binary files are easier and faster to transmit especially in web contexts [6].*
- *Also smaller files can be loaded into memory faster at runtime, which might lead to a slightly faster execution overall, but this is more of a speculation.*

2.2.3.F Modular

TODO:

- *Programs consist of smaller modules, which allows modules to be “transmitted, cached and consumed separably”*
- *Modules can also be combined/linked together at runtime*

2.2.3.G Other features

This section lists some other noteworthy features of Wasm. These are not directly related to this work, however they provide a better overview over Wasm’s potential use cases and applications.

- **Well-defined:** *TODO: Wasm is designed in such a way that it is “easy to reason about informally and formally” [6, sec. 1.1.1].*
- **Open:** *TODO*
- **Efficient:** *TODO: Wasm bytecode is efficient to read and parse, regardless of whether AOT or Just-in-time (JIT) compilation or interpretation is used at runtime [6].*

2.2 WebAssembly

- **Parallelizable:** *TODO: Working with Wasm bytecode should be easily parallelizable. This applies to all steps: decoding, validation, compilation. This property allows for a faster startup time.*
- **Streamable:** *TODO: This goal is especially important for the web. It should be possible to parse Wasm code while it is still being streamed/received. On the web data can be transferred in separate blocks called chunks. Wasm bytecode allows a Wasm runtime to decode, validate and compile a chunk before the full bytecode has arrived. This reduces startup time for Wasm applications especially on the web.*
- **Determinism:** *TODO: Indeterminism has only 3 sources: host functions, float NaNs, growing memory/tables*
- **Backwards-compatibility:** *TODO*
- **Migratability/Relocatability:** *TODO: Running Wasm instances can be serialized and migrated to another computer system. However for this to be easy the Wasm runtime should only execute Wasm code through interpretation.*

2.2.4 Challenges & Limitations

This section deals with common challenges and limitations of Wasm in non-web contexts.

TODO

2.2.5 WebAssembly System Interface (WASI)

TODO

2.2.6 WebAssembly Component Model

TODO

2.3 Plugin systems

TODO

3 Criteria for good plugin systems

To evaluate whether Wasm is a viable technology for versatile plugin systems, one must first understand what criteria make a plugin system good and versatile. This section will perform a technology comparison between several technologies and existing software projects. First, a set of criteria is defined. Then, appropriate technologies and software projects are selected to represent a spectrum of different plugin systems. Next, the technologies and projects are evaluated against the previously defined criteria. Finally, a technology comparison matrix is used to summarize and visualize the results.

3.1 Definition of criteria

In this section criteria for good plugin systems are defined. Each criterion will define a scale from 0 to 5, along with requirements for each score. This scale will be useful later to enable an objective evaluation and comparison of technologies and projects.

3.1.1 Performance

A computer's performance usually refers to the speed it is able to execute software at. For interactive computer systems one generally wants every piece of code to run as fast as possible to minimize its time on the CPU.

In the context of plugin systems, performance also refers to the speed at which software is executed. The three relevant software components necessary to define performance of a plugin system technology are the host system, the plugin system and the plugins managed and called by the plugin system. For this work we define performance as the property that describes how quickly a host system can temporarily transfer execution to a plugin system, which then loads and invokes a plugin's function.

While performance can be measured quantitatively through benchmarks, in practice this is quite hard for plugin systems. Plugin systems and their technologies often vary between host applications as they are by nature highly individual. To benchmark different plugin systems one would have to implement a variety of algorithms and scenarios for a variety of plugin systems. Then one could measure the time each plugin system and plugin takes to execute.

Due to time-constraints and the broad spectrum of knowledge about programming languages and host applications necessary, this work does not use quantitative benchmarks to measure performance. Instead performance is judged through educated

3.1 Definition of criteria

guesses based on benchmarks and comparisons already available for the technologies chosen and built upon by plugin systems.

0 – Very slow The transfer of execution to the plugin systems and invocation of a plugin is highly inefficient. Thus the plugin system is not viable for use within interactive software. Reasons for significant bottlenecks might include heavy serialization or expensive VM-based sandboxing.

1, – Slow Both the plugin system and plugins run very slowly. Transferring execution between the host system and a plugin is inefficient. Therefore this plugin system is also not recommended for use in interactive software, unless these inefficiencies can be somewhat mitigated, e.g. by offloading to other threads.

2, 3 – Acceptable The plugin system and plugins are not fast but their performance is acceptable for interactive software such as text editors. They can negatively impact the user experience by causing stuttering or slow loading times, but there are workarounds to minimize the impact of these problems.

4 – Fast Transferring execution to the plugin system and/or executing a plugin is fast, with only a small overhead, not noticeable by a user. While there is still a small overhead present, it is usually negligible in practice, except in scenarios with real-time requirements.

5 – Optimal Transferring execution and invoking a plugin is virtually instantaneous. There is no measurable overhead. All plugin code executes as fast as if it were implemented natively within the host system.

The scoring outline presented here is intentionally not very specific, without any hard lines between the different scores. It is meant to give only a rough guideline for evaluation, which then needs to be done very carefully on a case-by-case basis. For example, one could evaluate plugin systems based on whether plugins are compiled/interpreted or how large and thus slow plugins might be to load.

3.1.2 Plugin size

The plugin size property refers to the average size of a plugin for a specific plugin system technology. This property does not refer to the size of one specific plugin, but rather it is used to compare different plugin system technologies and how compact and small plugins for them are generally. The average plugin size may vary from technology

3.1 Definition of criteria

to technology due to factors such as static vs. dynamic linking of libraries or the size of the specific language's standard library.

The importance of plugin size depends on the specific use case and user requirements. For text editors specifically a smaller plugin size might result in faster startup times and less time spent downloading or updating the plugin. However reducing the overall impact a program has on the system's resource usage is generally preferred.

Also note that this section only refers to the plugin size and not the size of the entire plugin system. While the plugin system's size is also very important for the memory impact of the host application, it is harder to measure. This is due to the fact that plugin systems are usually very tightly coupled with the host application. To measure a plugin system's size, one would have to disable the plugin system of some host application, without breaking the host application itself. Only then would it be possible to compare system resource usages between the host application with and without a plugin system. Due to the high complexity, the plugin system's size will not be taken into account in this work.

The following scores will be used to evaluate a plugin's size. They are chosen specifically for plugin systems for text editors.

5 – Minimal(< 5KB) Plugin sizes are as minimal as they can possibly be. Plugins contain the minimal amount of program code necessary to achieve their desired functionality. The plugin code format is also made to be very space-efficient, which could be implemented for example through compression and hacks on the byte/bit levels.

4 – Negligible(< 500KB) Plugin sizes are so small that they are negligible in practice. There is no replication of similar information between multiple plugins such as statically-linked libraries.

3,2 – Moderate(< 50MB) Plugins are not very small, however their size is still quite manageable in the context of text editors. Examples could be plugin system technologies, that require a large fully-self contained runtime to be shipped with every plugin. Plugins might also have to contain all libraries and standard libraries that they depend on.

1 – Large(\leq 500MB) Plugins are unusually large specifically in the context of text editors. They are not as easy to manage and during runtime they might also have a non-negligible impact on the memory footprint of their host application.

3.1 Definition of criteria

0 – Very large(> 500MB) Plugins are very large. This may be due to the fact that their internal program logic requires very costly operations such as the virtualization of an entire environment, that must be completely self-contained in the plugin code.

3.1.3 Plugin isolation

Often times plugins contain foreign code. This is especially true for text editors, where plugins are often downloaded from a central registry, also known as plugin/extension marketplaces. This means that plugins downloaded from such sources are usually not validated and thus should be treated as foreign code. Even though there might be checks in place for malicious contents, foreign code should not be trusted to not access its host environment unless otherwise allowed.

For this work we define the property of plugin isolation to describe how isolated a plugin's environment is from its host environment. While there has to be some kind of interface between both environments to make plugins accessible and usable, this interface must not be considered when evaluating plugin isolation. Instead we define plugin isolation completely separated from the interface, meaning if an interface is unsafe by nature, plugin isolation is not automatically violated.

0 – No isolation, required elevated privileges Plugins are not only not isolated from the host application, they also require certain elevated privileges, which the host application usually does not require by itself.

Worst case: Elevated privilege access to current system.

1 – No isolation Plugins are not isolated from the host application. They inherit the host application's privileges without any attempt of the host plugin system to restrict these permissions.

Worst case: Plugins gain the same privileges as the host system, usually this means access to the current user's system and peripherals.

2, 3 – Restricted isolation Plugins are not isolated from the host application by design. Normally the plugin would inherit the host application's privileges, however the host application makes an attempt to restrict plugins from accessing certain critical functionalities. Some examples for restrictions on Linux systems could be allowing only a specific subset of syscalls through `seccomp(2)`⁶ or using

⁶<https://www.man7.org/linux/man-pages/man2/seccomp.2.html>

3.1 Definition of criteria

namespaces (7)⁷ to isolate and limit resources. However both of these examples do not use a sandboxing strategy for isolation.

Because these restrictions can come in a various shapes and forms each based on different technologies, during evaluation either 2 or 3 can be chosen as a score. *Worst case: Plugins have similar privileges to that of the host application, except for those specifically disallowed. However this restriction might be able to be circumvented.*

4 – Fully sandboxed, dynamic interface Plugins generally run completely isolated. However their interface is not statically defined, which can lead to vulnerabilities of the interface during runtime due to higher complexity and risk of bugs. Imagine a scenario where a host application exposes only a single interface for passing serialized messages back and forth with a plugin. Then the host application has to serialize and deserialize those messages during runtime. For complex systems, where advanced concepts such as additional shared memory between the host and plugins are used, this interface can become susceptible to logic bugs due to the dynamic interface.

Worst case: Access to parts of the host application not meant to be exposed due to a bug in the interface.

5 – Fully sandboxed, static interface The plugin runs fully sandboxed. It has no way of interacting with the host system, except for statically checked interfaces. Here statically checked interfaces refers to interfaces, that can be proven safe during compilation (or alternatively development) of the plugin system. One way to achieve this might be an interface definition in a common interface definition language. This restriction was chosen because it disallows plugin systems giving full access to parts of a host application without a proper interface definition.

Worst case: Indeterminable, a major bug in the sandboxing mechanism is required.

3.1.4 Plugin portability

Portability stems from the field of distributed systems. A high portability refers to software components, that can be moved from one distributed system A to another distributed system B without having to make any modifications[7]. This assumes that

⁷<https://www.man7.org/linux/man-pages/man7/namespaces.7.html>

3.1 Definition of criteria

both systems A and B share a common interface for interaction to the software component[7]. In the context of plugin systems for text editors, portability can be interpreted in one of two different ways:

1. Every individual plugin is seen as a software component. This plugin is portable, if it can be loaded into two instances of the same text editor running on different platforms.
2. The entire plugin system itself is seen as a modular software component of a text editor. It is portable if it can be integrated into different text editors and run across different platforms.

This work considers only the first scenario, in which portability refers to each individual plugin, because this scenario is less extensive and the portability of individual plugins is easier to measure. The following scores are used to measure plugin portability:

- 0 – Not portable** The plugin is not portable between different platforms. It is theoretically and practically impossible to run the plugin on different platforms.
- 1 – Theoretically portable** The plugin is theoretically portable between different platforms. In practice this might be very complex and costly, e.g. having to run each plugin in its own dedicated virtual machine.
- 2,3,4 – Portable with a runtime** The plugin is portable between different platforms, but it requires a runtime on the target platform. Because these runtimes can vary from one plugin system to another, a score range from 2 to 4 is specified here. During evaluation the specific runtime has to be analyzed regarding its complexity and impact on the host system. A more lightweight runtime could also enable higher portability of the plugin system itself as described in Section 3.1.4.
- 5 – Portable by design** The plugin is portable between different platforms without requiring a runtime on each host application. In practice this is very hard to achieve. Advanced technologies such as fat binaries, which are binaries that encapsulate compiled machine code for multiple different architectures, might be necessary.

Note that the most extreme scores 0 and 5 are very unlikely for any imaginable plugin system. 0 requires a plugin not to be portable at all, while 5 requires that a plugin is portable to different platforms and architectures which is near impossible to implement on a technical level.

3.1 Definition of criteria

3.1.5 Plugin language interoperability

Text editors and their plugin systems are highly individual software. Some users have personal preferences e.g. keybindings, macros or color schemes, while others may require tools such as language servers for semantic highlighting and navigation, or tools to compile and flash a piece of software onto an embedded device.

A lot of times plugin systems are used to overcome this challenge of high configurability. Major text editors and IDEs such as Neovim, VSCode, Emacs, IntelliJ or Eclipse provide plugin systems, some even in-built plugins. The advantage of implementing features as plugins is the reduced code complexity and size of the host application. Also providing a plugin systems with a publicly documented interface allows every user and developer to implement their own plugins for their own needs.

This property describes how interoperable plugins written in different languages are. In other words, the larger the set of languages is, in which a plugin can be written for a given plugin system technology, the better its plugin language interoperability is. The more languages are available, the less effort it requires for users to develop their own plugins without the need of learning a new (possibly domain-specific) programming language.

One could argue, that supporting a variety of different languages can result in higher interface complexity and less adaptability to new changes. Even though the complexity and adaptability of interfaces is another important property, which deserves its own rating, it will not be covered in this work.

0 – Domain-specific custom language Only a domain-specific language can be used to write plugins in. This language is specifically designed for given plugin system technology, which is why it is the least interoperable and might be the most unfamiliar and hardest for developers to learn and use.

1 – One language A single general purpose programming language is supported to write plugins in. There might be some developers who are already familiar with the language.

2 – Multiple languages A set of multiple languages is supported to write plugins in. The plugin system is able to abstract over multiple different plugin languages, so that the host application has only a single interface to communicate with plugins regardless of their specific language used.

3.1 Definition of criteria

- 3 – Build target for some languages** The plugin system supports a compilation target for plugins, that is targeted by multiple compilers for multiple programming languages. For example the Java bytecode is a compilation target for the Java, Kotlin and Scala compilers.
- 4 – Build target for a variety of languages** The plugin system supports a compilation target for plugins, that is targeted by a variety of different languages. This score differs from the previous score, that this score's compilation target is targeted by a considerably higher number of languages with more differences between them. Differences between languages could include dynamic vs. static typing, weak vs. strong typing, interpretation vs. just-in-time compilation vs. ahead-of-time compilation.
- 5 – Universal build target** The plugin system supports a compilation target to which most software can be compiled directly, or which can be embedded indirectly by a compiled runtime. For example all source code is eventually compiled to native ISA instructions specific to some hardware and platform. Thus it is also theoretically possible to package source code such as Python or JavaScript source code and combine it with their specific runtimes inside a native plugin.

3.2 A technology comparison between existing plugin systems

This section chooses appropriate plugin system technologies and software projects implementing a plugin system. Then the previously defined criteria will be evaluated for chosen technologies and projects.

3.2.1 Choice of technologies & projects

An important step during a technology comparison is the correct selection of technologies. When choosing technologies, one has to be careful to not introduce any bias towards certain technologies. This work mainly focuses on plugin systems for text editors, so text editor plugin systems will be the majority of technologies. However it could also be interesting to compare text editor plugin system technologies to plugin system technologies for other applications.

To make an appropriate decision for text editor projects, the popularity of different text editors will be as a metric for selection. As a source for the most popular text editors and integrated development environments (IDEs), the StackOverflow Developer

3.2 A technology comparison between existing plugin systems

Survey 2024 will be used[8]. The platform StackOverflow is known as a forum for developers helping each other by asking questions and exchanging information regarding various technical topics such as programming languages, certain APIs, technologies, etc. Annually it organizes a survey open to all developers regardless of their background to gather statistics about used programming languages, salaries of developers, used development tools, etc. In 2024 a total amount of ~65.000 developers took part in this survey, which is why it can be considered fairly representative. The four most popular text editors and integrated development environments (IDEs) according to the survey are the following. They are described and also considered for evaluation in this work:

- 1. Visual Studio Code** Visual Studio Code (abbreviated as VS Code) is a text editor designed for “coding”. It provides built-in basic features such as a terminal, version control or themes⁸. However it does not provide built-in integration for specific programming languages or technologies. In VS Code a lot of features are instead packaged as plugins (here: extensions) written in JavaScript which can be installed from a central marketplace. VS Code itself is also written in JavaScript. It is chosen for evaluation in this work, because it has successfully implemented a plugin system able to integrate a broad spectrum of plugins ranging from simple plugins for coloring brackets to highly complex plugins such as programming language integrations or plugins enabling interactive jupyter notebooks.
- 2. Visual Studio (not evaluated here)** Visual Studio is an IDE, that describes itself as “the most comprehensive IDE for .NET and C++ developers on Windows”⁹. It is not to be confused with Visual Studio Code, as Visual Studio provides more built-in features for developing, debugging, testing and collaboration between developers¹⁰. It also allows users to install plugins (here: extensions) written in C# from a central marketplace. However due to personal unfamiliarity with this IDE and the C# programming language together with the .NET ecosystem, this IDE will not be evaluated in this work.
- 3. IntelliJ-family** IntelliJ IDEA is an editor made by JetBrains for development of JVM-based languages such as Java or Kotlin¹¹. According to the StackOverflow developer survey, it is ranked as the third-most popular IDE.

⁸<https://code.visualstudio.com/>

⁹<https://visualstudio.microsoft.com/>

¹⁰<https://visualstudio.microsoft.com/vs/features/>

¹¹<https://www.jetbrains.com/idea/>

3.2 A technology comparison between existing plugin systems

IntelliJ's so called "Community Edition" is freely available and open-source, however JetBrains also develops proprietary alternatives for other use cases. In fact most of these alternatives such as CLion for C/C++ development or WebStorm for web development are also based on the IntelliJ framework. There are also third-party IDEs built upon the IntelliJ IDEA such as Android Studio for development of Android apps.

IntelliJ provides a plugin system where plugins can be written in Java or Kotlin, that can be used from all IDEs based on the IntelliJ framework. Thus the IntelliJ IDEA is not listed as a single IDE here, but rather a family of various IDEs all based on the same framework using the same plugin system technology.

- 4. Notepad++** Notepad++ is an open-source text editor focused around minimalism and efficient software¹². The editor itself is written in C++ with exclusive support for the Windows operating system. It provides a plugin system for loading plugins in the form of dynamically linked libraries (DLLs)¹³.

Besides these text editors and IDEs, this work also evaluates plugin system technologies for other application types. Multiple projects and technologies were considered, however due to their similarity to already chosen technologies, missing documentation and time-constraints for this work most of them were disregarded:

VST3 for music production During music production in a digital audio workstation, producers use plugins to simulate a variety of different audio effects, entire instruments or traditional analog devices such as compressors. One notable standard used for these kinds of plugins is the open and free Virtual Studio Technology (VST) 3 standard developed by Steinberg[9]. It specifies a standard for technology on how to implement a plugin system technology between a digital audio workstation (DAW) as the host application and a plugin which applies some effects or produces an audio signal. For example a DAW could send an audio stream to some plugin, then the plugin could apply an effect such as a simple equalizer on the audio signal and return the signal back to the DAW.

On the technical level VST3 plugins are DLLs on Windows, Mach-O Bundles on Mac and packages on Linux. The term package on Linux can have different meanings depending on the Linux distribution and package manager used.

¹²<https://notepad-plus-plus.org/>

¹³<https://npp-user-manual.org/docs/plugins/>

3.2 A technology comparison between existing plugin systems

The company behind the VST3 standard also provides a software development kit (SDK) for the C++ programming language and an API for the C programming language.

Microsoft Flight Simulator (not evaluated here) Microsoft Flight Simulator is a simulator for aircraft with focus on ultra-realism¹⁴. It allows for a wide variety of aircraft, airports and systems such as advanced flight information panels. This is achieved by providing multiple software development kits (SDKs) for plugins (so called add-ons). These SDKs support plugin languages such as C, C++, .NET languages, JavaScript or WebAssembly¹⁵.

However this software project will not be evaluated in this work. It's source code is not publicly available, which makes analysis of its architecture harder. Also a paid license is required to make an appropriate evaluation possible.

Eclipse(not evaluated here) Eclipse is an IDE used during software development for a variety of programming languages. It features a plugin system where plugins are written in Java and a central marketplace for installing plugins. However Eclipse is not chosen for evaluation, because its plugin system is too similar to IntelliJ's. Also IntelliJ is more popular[8] and thus required to support a greater variety of plugins.

3.2.2 Evaluations of technologies & projects

3.2.2.A Visual Studio Code

Performance In terms of performance Visual Studio Code performs reasonably well.

It builds on web technologies, specifically Electron which utilizes the Chromium browser engine together with Node.js. VS Code is written in JavaScript, which still imposes some limitations compared to native applications such as pauses due to garbage collection or the single-threaded nature of Node.js modules¹⁶. While JavaScript engines such as the V8 engine try to mitigate most issues, the performance of VS Code still remains slower and more memory-hungry than native alternative IDE and text editor applications written in C, C++, Rust, etc.

¹⁴<https://www.flightsimulator.com/>

¹⁵https://docs.flightsimulator.com/html/Programming_Tools/Programming_APIs.htm

¹⁶<https://www.electronjs.org/docs/latest/tutorial/multithreading#native-nodejs-modules>

3.2 A technology comparison between existing plugin systems

Even though the overhead imposed could be considered small and VS Code and its plugin system is quite fast for today's standards, its performance is still rated with a score of 3. This score is chosen, because VS Code performs below average but still acceptable in comparison to typical native applications.

Plugin size 3-5

- Most minimal plugins only contain JavaScript code, however extensions bundled for use of the text editor as a webapp include all dependencies
- *TODO: check sizes of common extensions*

Plugin isolation 4

- Sandboxed in theory
- no real permission system
- Interface is very dynamic

Plugin portability 4

- Plugins are very portable.
- JavaScript was designed with portability across various OS and platforms for all kinds of browser engines.
- However it still needs a JS engine/runtime

Plugin language interoperability 1

- only JS or typescript as its superset is supported
- although languages can be compiled to JS through technologies like asm.js, which might be deprecated in favour of Wasm?

3.2.2.B IntelliJ-family

Performance IntelliJ IDEs also perform reasonably well. Both the IDE, as well as its plugin system and plugins are written in languages targeting the Java Virtual Machine (JVM). That is, languages such as Java, Kotlin or Scala, which are optimized and compiled to Java bytecode by their respective compilers. This Java bytecode can then be executed by the JVM. During execution the JVM can apply additional optimizations and generate native machine code for execution for example through optimizing Just-in-time (JIT) compilation, only then when it is needed.

Even though JVMs are highly complex systems with loads of mechanisms for optimization, they still introduce some overhead. Also they are responsible for

3.2 A technology comparison between existing plugin systems

holding the state of programs during runtime and for garbage collecting no longer used object instances.

Thus the performance of plugin systems and plugins based on the JVM, will not be able to outperform native non garbage-collected approaches. This plugin system technology, as it is used in IntelliJ-based IDEs is rated with an average performance score of 3.

Plugin size 2-3

- Java links dependencies, except for the standard library statically
- *TODO: check sizes of common plugins*

Plugin isolation 2

- All plugins run in the same JVM only with different classloaders. This makes isolation impossible

Plugin portability 4

- Java was designed with portability across devices of all kinds in mind.

Plugin language interoperability 3

- Plugins are executed by a JVM.
- SDKs are available for Java & Kotlin, however all other JVM-based languages could also be used theoretically

3.2.2.C Notepad++

Performance Notepad++ itself is written in C++ and compiled to native machine code for the Windows operating system exclusively. It tries to maximize efficiency and minimize the impact on the system it is running on. Its plugin system is based on compiled dynamically linked libraries (DLLs). A plugin developer compiles their plugin, which can be written in any arbitrary language to a DLL, which is essentially a library containing machine code along with symbols defining imports and exports of that library. Notepad's plugin system can then load these libraries at runtime via the LoadLibrary Win32-API call and execute arbitrary functions exported from the plugin.

This is the fastest way for how some host application can embed plugins. It relies only on the operating system for loading already compiled machine code at runtime. There is no additional overhead except the time needed for loading the DLL itself into memory.

3.2 A technology comparison between existing plugin systems

Thus Notepad++'s plugin system is evaluated at a score of 5 for its optimal performance.

Plugin size 3-4

- *TODO: check sizes*

Plugin isolation 1

- Plugins are loaded as dynamic libraries at runtime without any isolation.

Plugin portability 1

- Notepad++ itself does not support multiple platforms.
- Thus there is also no need for plugins to be portable across different platforms

Plugin language interoperability 5

- native code is a common build target for all languages

3.2.2.D VST3

Performance The VST3 standard for plugins creating and processing audio relies on native compiled machine code just like Notepad++. However it's file format also accommodates for the fact that plugins might be run on more than one platform. Thus the VST3 format allows for embedding of DLLs for Windows plugins, Macho-O bundles for MacOS plugins or Packages for Linux plugins.

During runtime a host application has to check whether a VST3 plugin contains machine code compiled for the current architecture. Then it is able to link the machine code at runtime through the operating system just like Notepad++ does.

While there is a small overhead of checking if the targeted platform of a plugin is correct for loading a plugin, there is no overhead during execution of plugin code. Thus the VST3 technology also receives a score of 5.

Plugin size 3-4

- Same as Notepad++ probably
- However here plugins contain more data such as images for user interfaces on average? *TODO: is this true?*

Plugin isolation 1

- No isolation

Plugin portability 1-2

3.2 A technology comparison between existing plugin systems

	<i>Performance</i>	<i>Plugin size</i>	<i>Plugin isolation</i>	<i>Plugin portability</i>	<i>Plugin language interoperability</i>
Visual Studio Code	3	4	4	4	1
IntelliJ-family	3	3	2	4	3
Notepad++	5	4	1	1	5
VST3	5	3	1	1	5

Table 1: Technology comparison matrix for selected technologies and software projects

- Compiled plugins are not portable as they are compiled for a specific architecture and platform. e.g. Windows (DLLs), MacOS (Mach-O Bundle) or Linux (package)
- However the source code of plugins can be reused across multiple platforms according to the VST3 documentation.

Plugin language interoperability 5

- native code is a common build target for all languages

3.2.3 Summary

TODO: Present findings in a table TODO: What could have been done better?

- *Complexity and adaptability of the interface*

TODO: Which other technologies and criteria might also be interesting? Which ones were left out?

4 WebAssembly for plugin systems

TODO: Present basic idea of running Wasm code for each plugin inside a Wasm runtime

4.1 Overview of basic plugin system architecture

4.2 Evaluation of requirements

4.3 Evaluation of interface-specific requirements

It is not possible to evaluate Wasm as a technology, because it only provides fundamental technology for executing WebAssembly program code. In practice WebAssembly is often combined with other technologies that build upon Wasm's basic constructs and allow for more complex systems such as type systems or interface definitions.

WebAssembly as a technology is often not restrictive enough and thus the decision of whether a requirement is fulfilled often comes down to the host- and plugin-language and whether a common interface definition between them exists.

To illustrate this point, consider a scenario in which a host system is written in JavaScript. When this host system wants to call a Wasm function it serializes the arguments, which might consist of complex JavaScript types, to JSON strings. Then it passes these JSON strings to the Wasm plugin. A plugin written in JavaScript itself will be able to easily parse the JSON string given to it, however a plugin written in C first has to get a system in place to parse and convert JavaScript types to equivalent C types.

TODO: Wasm interfacing is still an unsolved problem. There are many different solutions, of which some will be evaluated separately here

4.3.1 Without a standardized interface

TODO

4.3.2 WebAssembly Component Model

TODO

4.3 Evaluation of interface-specific requirements

4.3.3 WebAssembly Component Model + WebAssembly System Interface (WASI)

TODO

4.3.4 Custom serialization format (JSON, XML, Protobuf)

TODO

4.4 Plugin systems already using WebAssembly today

Zed text editor with Wasm plugin system, no windows support

- Wasm, but official interface only for Rust plugins?
- WASM mit kompletter WIT Schnittstellendefinition (Generisches Typ/Funktions-basiertes System)pro Version
- Fokus auf Isolation von WASM Code
- shim library nur für Rust vorhanden: `zed_extension_api`
<https://zed.dev/docs/extensions/developing-extensions> https://github.com/zed-industries/zed/blob/94faf9dd56c494d369513e885fe1e08a95256bd3/crates/extension_api/wit/since_v0.2.0/http-client.wit

Zellij terminal multiplexer, has a Wasm plugin system.

- Wasm, but official interface only for Rust plugins?
- Zellij is a terminal workspace (similar to a terminal multiplexer). It is used to manage and organize many different terminal instances inside one terminal emulator process. Similar commonly known terminal multiplexers are Tmux, xterm or the Windows Terminal.
- plugin system to allow users to add new features
 - plugin system is not very mature <https://zellij.dev/documentation/plugin-system-status>
 - Wasm for plugins, however only Rust is supported
 - Permission system

Extism generic Wasm plugin system library usable in many different languages

- Extism is a cross-language framework for embedding WebAssembly code into a project.
- It is originally written in Rust and provides bindings (Host SDKs) and shims (Plugin Development Kits: PDKs) for many different languages.
- Docs: <https://extism.org/docs/overview>

4.4 Plugin systems already using WebAssembly today

- Github: <https://github.com/extism/extism>

go-plugin *TODO*

4.5 Summarized evaluation for WebAssembly

TODO: Show all WebAssembly configuration in a matrix with all criteria as columns

5 Proof of concept: Implementing a WebAssembly plugin system for a text editor

TODO: Provide context of helix text editor

5.1 Requirements

TODO: Weight the requirements for plugin systems and optionally add new requirements for this project

5.1.1 Functional requirements

5.1.2 Non-functional requirements

5.2 System Architecture

TODO: Choose appropriate Wasm technologies based on the previous findings and document how they are used together to build a working system

5.3 Implementation

TODO: Give brief overview over code structure TODO: Technical challenges and how they were addressed TODO: What optimizations were made?

5.4 Evaluation & Results

TODO: Reevaluate the key requirements for plugin systems TODO: Provide measurements for memory/performance impact (there are no real reference points) TODO: Summarize findings & challenges

Standard plugin definieren (z.b. textsuche für performance) Graph mit x geladenen
Standard plugins für memory/performance

6 Results & Discussion (2 pages)

- some of wasms strenghts were not represented during the technology comparison.
namely native-like safe interfacing

7 Outlook

Bibliography

- [1] A. S. Tanenbaum and T. Austin, *Structured Computer Organization*, 6th ed. Pearson Education, Inc., 2025.
- [2] A. Haas *et al.*, “Bringing the Web up to Speed with WebAssembly,” in *ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI) 2017*, Association for Computing Machinery, Jun. 2017, pp. 185–200. doi: 10.1145/3062341.3062363.
- [3] W. Zaeske, S. Friedrich, T. Schubert, and U. Durak, “WebAssembly in Avionics: Decoupling Software from Hardware,” in *IEEE/AIAA 42nd Digital Avionics Systems Conference (DASC)*, Nov. 2023. doi: 10.1109/dasc58513.2023.10311207.
- [4] M. N. Hoque and K. A. Harras, “WebAssembly for Edge Computing: Potential and Challenges,” in *IEEE Communications Standards Magazine*, IEEE, 2022, pp. 68–73. doi: 10.1109/MCOMSTD.0001.2000068.
- [5] S. Wallentowitz, B. Kersting, and D. M. Dumitriu, “Potential of WebAssembly for Embedded Systems,” in *11th Mediterranean Conference on Embedded Computing (MECO)*, IEEE, Jun. 2022. doi: 10.1109/meco55406.2022.9797106.
- [6] WebAssembly Community Group and A. Rossberg, “WebAssembly Core Specification (Release 2.0, Draft 2025-01-28).” Accessed: Mar. 27, 2025. [Online]. Available: <https://webassembly.github.io/spec/core>
- [7] M. van Steen and A. S. Tanenbaum, *Distributed Systems: Principles and Paradigms*, Edition 4, version 03. Maarten van Steen, 2025. Accessed: Mar. 27, 2025. [Online]. Available: <https://www.distributed-systems.net/index.php/books/ds4/>
- [8] “StackOverflow Developer Survey 2024.” Accessed: Apr. 08, 2025. [Online]. Available: <https://survey.stackoverflow.co/2024/>
- [9] “VST 3 Developer Portal.” Accessed: Apr. 09, 2025. [Online]. Available: https://steinbergmedia.github.io/vst3_dev_portal