

[Goto table of contents](#)



Exploring WebAssembly for versatile plugin systems through the example of a text editor

im Studiengang TINF22IT1

an der *Duale Hochschule Baden-Württemberg Mannheim*

von

Name, Vorname: Hartung, Florian

Abgabedatum: 15.04.2025

Bearbeitungszeitraum: 15.10.2024 - 15.04.2024

Matrikelnummer, Kurs: 6622800, TINF22IT1

Wiss. Betreuer*in der Dualen Hochschule: Gerhards, Holger, Prof. Dr.

Erklärung zur Eigenleistung

Ich versichere hiermit, dass ich meine Projektarbeit mit dem

THEMA

Exploring WebAssembly for versatile plugin systems through the example of a text editor

selbstständig verfasst und keine anderen als die angegebenen Quellen und Hilfsmittel benutzt habe. Ich versichere zudem, dass die eingereichte elektronische Fassung mit der gedruckten Fassung übereinstimmt.

Ort, Datum

Unterschrift

Abstract

Plugin systems are a software architecture pattern for making host applications extensible without modifying the host application itself. They are used in various applications such as real-time audio processing or computer games, but they play a particularly important role in text editors. Text editors and IDEs are applications used primarily by developers for software development. Depending on the required technologies and highly individual developer preferences, text editors must be able to adapt to a variety of different use cases. Such use cases may include language support, complex keybindings or development on a remote server.

WebAssembly (Wasm) is a relatively new technology first released in March 2017. It was originally designed as a fast, safe and portable compilation target for higher level languages to enable fast code execution on the web. However, it is designed with no assumptions about its execution environment and its properties make it an interesting technology for non-web contexts such as avionics, distributed computing, embedded devices or automotive as well.

This work explores Wasm as a technology for implementing a safe and fast plugin system with portable plugins compiled from higher level languages, such as C, C++, Rust, Python or JavaScript. The question is whether Wasm, as a new and versatile technology, can outperform existing plugin system technologies in terms of performance, security, compatibility and ease of development.

To do this, a technology comparison is conducted between selected existing plugin systems to analyze the state of the art of plugin system technologies. WebAssembly will then be evaluated against the same criteria to provide an objective comparison with existing technologies. Finally, as a proof of concept, a plugin system based on WebAssembly will be developed for an existing text editor project to provide a better insight into the current state of the rapidly evolving Wasm ecosystem.

TODO: present results

TODO: discuss results and impact of this work briefly

Zusammenfassung

TODO: Abstract übersetzen

Contents

1 Introduction	1
1.1 Motivation & problem statement	1
1.2 Research question	2
1.3 Methodology	2
2 Fundamentals	2
2.1 Instruction set architectures	3
2.2 WebAssembly	3
2.2.1 Overview	3
2.2.2 Execution model and lifecycle	5
2.2.3 Design goals	7
2.2.4 Challenges & Limitations	11
2.2.5 WebAssembly Component Model	13
2.2.6 WebAssembly System Interface (WASI)	15
2.3 Plugin systems	16
3 Technology comparison of existing plugin systems	17
3.1 Definition of criteria	18
3.1.1 Performance	18
3.1.2 Plugin size	19
3.1.3 Plugin isolation	21
3.1.4 Plugin portability	22
3.1.5 Plugin language interoperability	23
3.2 Choice of technologies & projects	25
3.3 Evaluations of technologies & projects	28
3.3.1 Visual Studio Code	28
3.3.2 IntelliJ-family	31
3.3.3 Notepad++	32
3.3.4 VST3	34
3.4 Results of the technology comparison	36

4 WebAssembly for plugin systems	37
4.1 Evaluation of criteria	39
4.2 Evaluation of WebAssembly for plugin systems against previous technologies	41
4.3 Related projects using WebAssembly plugin systems	42
4.4 Results	43
5 Proof of concept: Implementing a WebAssembly plugin system for a text editor	44
5.1 Requirements	45
5.1.1 Functional	45
5.1.2 Non-functional	46
5.1.3 Non-requirements	46
5.2 System Architecture	46
5.3 Implementation	48
5.4 Verification and validation	50
5.4.1 Verification	50
5.4.2 Validation	51
5.4.3 Conclusion	52
6 Results & discussion	52
7 Outlook	53
Bibliography	55

List of figures

Figure 1: A multilevel computer system running WebAssembly code	4
Figure 2: Flowchart of the lifecycle of a WebAssembly module	5
Figure 3: A plugin system inside of a host application	16
Figure 4: Plugin size distributions of selected Visual Studio Code plugins	28
Figure 5: Plugin size distributions of selected Notepad++ plugins	33
Figure 6: A WebAssembly plugin system inside of a host application	38
Figure 7: Dataflow for loading Wasm plugins	46

List of tables

Table 1: Technology comparison matrix of existing technologies & projects	36
Table 2: Final technology comparison matrix for existing technologies & WebAssembly	41

List of code listings

Listing 1: C and Rust functions and their WebAssembly equivalents	11
Listing 2: An exemplary WIT definition	13
Listing 3: WIT definition with all interfaces for the developed WebAssembly plugin system....	49
Listing 4: WIT definition for a single example Wasm plugin	49

1 Introduction

Web applications are becoming more and more popular and complex requiring the use of compute-intensive algorithms for applications such as 3D visualizations or interactive games[1]. Thus the need for faster code execution that cannot be provided by JavaScript grows. Traditionally JavaScript is the only language allowing client side code execution, however it was never designed for these kinds of high-performance tasks.

WebAssembly (Wasm) is a modern web technology developed through collaboration of the four major web browsers providing performance, safety and portability. It is used by popular frameworks such as Qt for building web interfaces powered by C++ code [2] or companies such as Ebay to optimize a slow barcode scanner algorithms[3].

TODO: Explain why Wasm is attractive for non-web contexts

TODO: Name some non-web contexts, where Wasm is already being used and actively being researched on

TODO: Plugin systems on the other hand are software architectures widely used, especially in text editors. TODO: Explain the advantages of plugin systems and why they are so important for text editors TODO: Name some problems plugin system might have

TODO: Wasm looks promising as a technology for plugin systems

1.1 Motivation & problem statement

TODO: Motivate in two parts:

- 1. Plugin systems are suitable architecture for highly individual software like text editors*
- 2. WASM is relatively new bytecode that provides speed, safety, portability by default*

TODO: Problem statement:

- Plugin systems suffer from many issues: Security, interoperability, Portability, (Developer experience?)*
- WASM's use cases are very limited -> it is still very new and evolving*
- Can both be combined for better plugin systems?*

1.1 Motivation & problem statement

1.2 Research question

Is WebAssembly the best technology choice for designing versatile plugin systems especially for text editors?

1.3 Methodology

TODO: How this work is structured:

- *Technology comparison*
- *Additional evaluation and comparison of Wasm*
- *Development of a proof of concept*

2 Fundamentals

This section introduces theoretical and technical fundamentals used in this work. First it covers the definition of an instruction set architecture. Then it gives an overview over WebAssembly, its features, challenges, limitations and extensions. Finally, plugin systems are explained as a software architecture model.

2.1 Instruction set architectures

In his book “Structured Computer Organization” Tanenbaum defines an instruction set architecture (ISA) as a level in a multilayered computer system[4, sec. 1.1.2]. The ISA level defines a machine language with a fixed set of instructions. According to Tanenbaum the ISA level then acts as a common interface between the hardware and software. This allows software in the form of ISA instructions to manipulate the hardware. Software written in a higher level machine language (Assembly, C, Java, ...) can not be executed directly by the hardware. Instead higher level machine codes are compiled to ISA machine code or interpreted by a program, that is present in ISA machine code itself [4, sec. 1.1.2].

2.2 WebAssembly

2.2.1 Overview

WebAssembly (Wasm) is a stack-based ISA for a portable, efficient and safe code format. Originally it was designed by engineers from the four major vendors to enable high-performance code execution on the web [1]. However it is also becoming increasingly interesting for researchers and developers in non-web contexts. Some examples are avionics for Wasm’s safe and deterministic execution [5], distributed computing for its portability and migratability [6] or embedded systems for its portability and safety [7].

What is special about Wasm is that it is a *virtual* ISA [8, sec. 1.1.2]. There is no agreed-upon definition for a virtual ISA, however the term *virtual* can be assumed to refer to an ISA that is running in a virtualized environment on a higher level in a multilevel computer¹. We call this virtualized environment the **host environment**

¹Projects that try to execute Wasm directly exist. One example is the discontinued wasmachine project, which tried executing Wasm on FPGAs: <https://github.com/piranna/wasmachine>

2.2 WebAssembly




Figure 1: A multilevel computer system running Wasm code. Based on [4, fig. 1-2].

(used by the specification[8, sec. 1.2.1]) or the **WebAssembly runtime** (used by most technical documentation).

If one considers a system running Wasm code as a multi-level computer system, the Wasm runtime can be modeled as a separate layer. Figure 1 shows a multi-level computer system based on Tanenbaum’s model [4, sec. 1.1.2]. Here each level is executed by logic implemented in the next lower level either through compilation or interpretation. The digital logic level itself only exists in the form of individual gates, consisting of transistors and tracks on the processors’ chip. This level runs the next microarchitecture and ISA levels, which are also often implemented directly in hardware. The ISA level then provides a fixed set of instructions for higher levels to use. Operating systems build on top of this and provide another level for user space programs, which exist on the assembly language level. Then there are problem-oriented languages such as C, C++ or Rust, which are specifically made for humans to write code in[4, sec. 1.1.2].

One program written in a problem-oriented language is the Wasm runtime, which itself is a layer here. Its task is to interpret or compile higher-level Wasm code to lower-level problem-oriented or even the assembly language level. However for this work all layers starting with the Wasm runtime level until the digital logic level can be seen as a single hardware specific layer called the *Wasm runtime environment*.

WebAssembly code can be written by hand, just as one could write traditional ISA instructions (think of writing x86 machine code) by hand. Even though it is possible, it would not be efficient to write useful software systems in Wasm because the language is too simple. For example it provides only a handful of types: Signed/unsigned integers,

2.2 WebAssembly

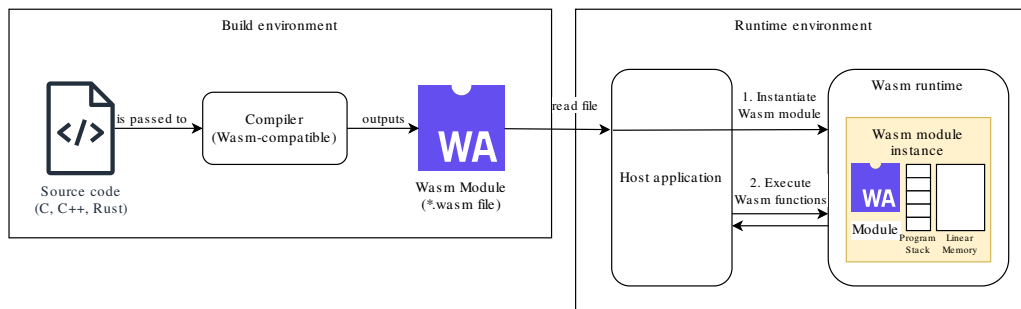


Figure 2: Flowchart for the creation of a Wasm module from a higher-level language and execution through a host application.

floating point numbers, a 128-bit vectorized number and references to functions/ host objects.

WebAssembly also does not provide any way to specify memory layouts as it can be done in higher-level languages with structs, classes, records, etc. Instead it provides most basic features and instructions, which exist on almost all modern computer architectures like integer and floating point arithmetic, memory operations or simple control flow constructs. This is by design, as WebAssembly is more of a compilation target for higher level languages[8]. Those higher level languages can then build upon Wasm’s basic types and instructions and implement their own abstractions like memory layouts or control flow constructs on top. This is analogous to the non-virtual ISA machine code in a conventional computer, which also acts as a compilation target for most low-level languages such as Assembly, C, or Rust. Nowadays there are compilers for most popular languages already. Some of which are `clang` for C/C++, `rustc` for Rust or the official Zig and Haskell compilers. However most compilers are still being actively worked on and improved over time to support the latest Wasm proposals.

2.2.2 Execution model and lifecycle

Figure 2 shows the different stages a Wasm module goes through in its lifecycle. Its lifecycle starts with a developer writing source code. This source code can be written in an arbitrary programming language such as C, C++ or Rust. These languages are often used because they have compilers that support Wasm. The compiler then compiles this source code to a *Wasm module*. This step requires a compiler that support Wasm as a

2.2 WebAssembly

compilation target. Many modern compilers such as `clang` or `rustc` use the LLVM² project for optimization and code generation. These compilers compile source code to LLVM's intermediate representation (LLVM IR) and pass this to LLVM. LLVM can then perform optimizations and compile this IR to any compilation target, which can be selected by choosing a LLVM backend. One such LLVM backend targets Wasm. This way the Wasm-specific compiler logic can be implemented once within LLVM, and any compiler using LLVM can then target Wasm with minimal additional effort.

The compiled Wasm module exists in the form of a `.wasm` file. It is fully self-contained and unlike binaries it cannot depend on any dynamic libraries at runtime. It consists of different ordered sections, each with their own purpose: Some example sections contain function signatures, data segments, import- and export definition or actual Wasm instructions for each function.

The previously generated Wasm module can then be transferred to any target device or platform providing a Wasm runtime. Wasm runtimes are usually hardware-specific and not portable, unlike the Wasm module. The Wasm runtime is able to parse the received Wasm module file, instantiate a *Wasm instance* from it and provide an Application Programming Interface (API) for interaction with the Wasm instance. APIs can differ from one Wasm runtime to another. Some runtimes exist as standalone programs that can run Wasm modules comparable to how native binaries can be executed. Others like the one shown in [Figure 2](#) are in the form of libraries, that can only be used from a host application to embed a Wasm runtime into them. These Wasm runtime libraries often provide common operations to the host application like calling Wasm functions, reading and writing operations for Wasm memories, linking mechanisms between Wasm modules, exposing host-defined functions for Wasm instances to call, etc.

In a web context a server might provide this Wasm module to the client's browser, which usually comes with a Wasm runtime. Detailed information on the current status of Wasm support for internet browsers can be viewed at <https://caniuse.com/wasm>. A concrete example is Ebay using Wasm for their barcode scanner algorithm to achieve higher performance³.

²According to its official website the acronym *LLVM* was once short for *Low Level Virtual Machine*, however now it has “grown to be an umbrella project” referred to simply as the LLVM project. <https://llvm.org/>

³<https://innovation.ebayinc.com/tech/engineering/webassembly-at-ebay-a-real-world-use-case/>

2.2 WebAssembly

For distributed computing and especially serverless functions a compiled Wasm module could be distributed among multiple different nodes regardless of their platform and architecture while still providing good performance. This is what the projects WasmEdge or Fermyon (see <https://wasmedge.org/> and <https://www.fermyon.com/>) are trying to achieve.

2.2.3 Design goals

Wasm was designed with certain design goals in mind. This section presents the design goals relevant for this work according to the official Wasm specification [8, sec. 1.1.1]. Each design goal is accompanied by related information from various papers and articles for a deeper understanding of each goal.

2.2.3.A Fast

Wasm is designed to be fast both during startup and execution [8, sec. 1.1.1]. Startup time is mostly optimized through the structure of the Wasm bytecode format, which is optimized for fast parsing and compiling of Wasm code. To further elaborate, a Wasm module in its binary format consists of 11 different sections. These sections must only occur in a very specific order to allow Wasm runtimes to parse and/or skip them efficiently. Another example is the absence of backwards jump (except for the loop statement), which allows the use of faster one-pass compilers.

During runtime Wasm's execution time varies. The original design paper for Wasm found Wasm to be at most two times as slow as native code and around 30% of PolyBenchC's benchmarks to close to native code execution time [1, p. 197]. Another performance analysis by Jangda et al. measures peak slowdowns of 2.5x (Chrome) and 1.45x (Firefox) for running Wasm code in different browsers and comparing their execution times to native execution[9].

An additional property of Wasm bytecode is that it can be either compiled, interpreted or just-in-time compiled by a runtime. This provides users of Wasm with flexibility for performance optimizations. Some use cases require the fast startup time, determinism and relocatability (see Section 2.2.3.F) only an interpreter can provide. Others require high performance at the cost of flexibility of a compiler or just-in-time compiler.

2.2 WebAssembly

2.2.3.B Safe

Wasm’s execution is safe and sandboxed by default. All operations operate on a stack, which contains values, labels as references for branch instructions and activations, which are similar to function call frames. The Wasm stack differs from a stack used in most native architectures in two ways: Wasm code can only access the top-most value of the stack, and the stack is type-checked by the Wasm runtime during a validation phase[8].

Wasm modules also have their own linear memory. This linear memory is completely isolated from the stack and can only be accessed using special instructions and bounds-checked indices instead of pointers. While host applications can always read from and write to Wasm linear memory, Wasm instances cannot access host memory in the opposite direction. Instead, Wasm instances can only interact with the host through functions explicitly exposed by the host[8].

While the Wasm execution format is sandboxed, it is still important that the implementing Wasm runtime does not introduce any sandbox escapes due to bugs, etc. In their article “Provably-Safe Multilingual Software Sandboxing using WebAssembly” Bosamiya et al. have acknowledge Wasm as a safe sandboxed execution format and implemented a provably safe and verifiable Wasm runtime[10]. This shows that implementation of a fully sandboxed and safe runtime is possible not only in theory but also in practice.

Because Wasm is able to provide safety for untrusted code execution, it is currently gaining interest in certain safety-critical non-web contexts, such as avionics[5] or automotive industries (see <https://oxidos.io/>).

2.2.3.C Portable

Originally Wasm was designed for fast and safe code execution on a client’s web browser[1]. This means that a variety of web browsers, all running on different platforms and architectures. Various devices such as desktop computers, mobile devices and even embedded devices have to be able to run Wasm code.

Thus Wasm was designed to be able to be portable for various target platforms and architectures[8]. It does this by defining only the most basic types and instructions necessary to act as a compilation target. The basic types and instructions chosen exist on most platforms and architectures today. For example Wasm defines ~172 different

2.2 WebAssembly

instructions and ~8 types[8, sec. 4.2.1]. In comparison, the x64 architecture defines ca. 1500 - 6000 instructions[10].

2.2.3.D Independence of language

Wasm is designed to be a compilation target for higher level languages. In particular it is not designed for a specific language, programming model or object model[8, sec. 1.1.1]. Currently Wasm can already be used as a compilation target by various languages such as C, C++, Rust, Haskell or Ada.

2.2.3.E Compact

The representation of the Wasm bytecode format is designed to be compact. Especially on the web file size is very important to minimize loading times and achieve better user experiences. Also smaller files are faster to load into memory, which might lead to a slight increase in performance.

Research shows that the compactness of Wasm bytecode in non-web contexts is close to native code with around 246% the size of natively compiled x86_64 machine code for the PolyBenchC benchmarks[11]. While Wasm should optimally come as close to native code compactness as possible, this is close to impossible as native code is able to provide a larger variety of instructions exposing hardware-specific behavior.

As a concrete examples for how Wasm bytecode is compacted, variable integers are explained in the following: Often times, the top-most bytes of integers are 0. To reduce the size of integers, their bit information is not stored as a fixed amount of bytes (e.g. 4 bytes for a 32-bit integer) anymore. Instead only 7 bits per byte are used to store the integer's bits, while the 8th bit is reserved for a flag that indicates whether the next byte also belongs to the integer currently being encoded. This allows to store 32-bit unsigned integers in the interval $[0; 127]$ inside one byte, $[128; 16383]$ inside two bytes, etc., while the maximum possible integer value of $2^{32} - 1$ now requires 5 bytes for storage. This design is very similar to how Unicode encodes multi-byte characters.

2.2.3.F Other features

This section lists some other noteworthy features and design goals of Wasm. These are not directly related to this work, but instead provide a better overview over Wasm's potential use cases and applications.

2.2 WebAssembly

- **Determinism:** The majority of Wasm code is deterministic and undefined behavior is prevented through validation before execution of Wasm code happens. There are only three possible sources of indeterminism:
 - Functions exposed by the host application and imported into a Wasm instance can cause side-effects and introduce indeterminism[\[8, sec. 4.4.10\]](#).
 - NaN float values and operations on them are not defined deterministically, because there exists a variety of different NaN values[\[8, sec. 2.2.3\]](#).
 - Wasm instances can grow their memory or tables through the `memory.grow` and `table.grow` instructions. Wasm runtimes are able to make these instructions fail, either due to missing resources, because of resource limitations or any other reason without determinism[\[8, sec. 4.4.6, sec. 4.4.7\]](#).
- **Parallelizable:** Wasm bytecode is designed for operations on it to be easy to parallelize[\[8, sec. 1.1.1\]](#). This applies to decoding, validation and compilation of Wasm bytecode and it may allow for faster startup times of Wasm instances.
- **Streamable:** With Wasm being made for the web, the streamable property is quite unique from other ISAs. This property means that Wasm bytecode is able to be read and consumed by a Wasm runtime, before the entire data has been seen[\[8, sec. 1.1.1\]](#). On the web this means, that the Wasm runtime inside a client's browser can start parsing and validating a Wasm module while it is being received over the internet. This reduces startup time for Wasm instances especially on the web, but it may also be relevant for non-web contexts, e.g. where large Wasm modules are being read from slow hard disk drives.
- **Modular:** Wasm applications consist of smaller modules, so called Wasm modules. They are self-contained, which allows them to be “transmitted, cached and consumed separately”[\[8, sec. 1.1.1\]](#). During runtime most Wasm runtimes also allow multiple Wasm modules to be linked together, as long as all imports required by those Wasm modules are fulfilled.
- **Well-defined:** Wasm code is designed to be easy to understand and “reason about informally and formally”[\[8, sec. 1.1.1\]](#) without loopholes or undefined behavior.
- **Relocatability:** Wasm modules by themselves are clearly defined at runtime consisting of their bytecode, a linear memory, reference tables, etc. Thus it is possible to halt execution of Wasm instances, serialize the entire runtime state and relocate it to another computer system for continuation of execution in theory. This could enable

2.2 WebAssembly

```
char* get_c_string() {  
    ...  
}  
                                     // C  
                                     (type (result i32))  
  
fn get_rust_string() -> String {  
    ...  
}  
                                     // Rust  
                                     (type (result i32) (result i32))
```

(a) Functions signatures in the C and Rust programming languages. (b) WebAssembly functions in its text format for the C and Rust programs.

Listing 1: Function signatures in the C and Rust programming languages and their compiled type signatures in the WebAssembly text format after compilation. Both functions take no arguments and return a string.

new methods for debugging, where traditionally a runtime state and logs are dumped for later analysis by a developer. Now an entire snapshot of the runtime state could be created, that can be rerun and analyzed at a later time.

2.2.4 Challenges & Limitations

This section deals with common challenges and limitations of Wasm especially important in non-web contexts.

Wasm provides only a very basic type system with 8 types: 32-bit and 64-bit integers and floats, a 128-bit vector type and three types of references to functions, host objects and the null reference. Its type system does not contain any methods to combine multiple data types into a new one, such as structs or classes which are common in most programming languages. This is intentional, so that Wasm can be a universal compilation target independent of the higher-level language used (see [Section 2.2.3.D](#)), because not all languages share the same data constructs. However this basic type system also has its downsides. Because Wasm delegates the responsibility of managing data layouts and representations to compilers, interfaces are now depending on the specific compiler used to create the Wasm module. Consider as an example the different strings representations in C and Rust. Traditionally strings in C exist as pointers to null-terminated byte sequences of ASCII characters. In Rust strings use the Unicode format by default. Rust strings still store a pointer to byte sequence, however characters are encoded, allowing for multi-byte characters for example. Also Rust strings do not use

2.2 WebAssembly

null-terminators, and instead store the length of the string alongside the pointer in a `String` struct.

Now if one were to compile a C program exposing a function returning a string and a Rust program exposing a function that also returns a string, these would compile to different Wasm function types.

[Listing 1a](#) shows two function signatures in the C and Rust programming languages for functions returning a string. When these functions are compiled to WebAssembly, they result in two different function signatures shown in [Listing 1b](#). The compiled C function returns a pointer in the form of an integer to a null-terminated string. The compiled Rust function returns two integers, one as the pointer to the bytes of a Unicode string and the other as the length of the string in bytes.

This issue exists for the various types commonly found in programming languages, for example consider the data layout of enums, tagged unions, futures, options or results. There are multiple solutions, for how Wasm interfaces can be standardized across programming languages:

1. A common serialization format such as JSON, XML or a binary format such as CBOR can be used. Before the host invokes a Wasm function, it has to serialize all arguments into a large byte array. Then this byte array is passed indirectly as an argument, by writing it into the Wasm instance's linear memory and passing an index to the data to the function. Internally the function first has to deserialize the data into its original form, before it can execute its program logic. At the end of the function, the same procedure is repeated to pass return values back to the host. For this approach a library such as Protobuf (see <https://protobuf.dev/>) could also be used. It provides serialization and deserialization logic for a variety of programming languages. However this approach might introduce both a performance and memory overhead, due to the heavy serialization and deserialization necessary.
2. Another approach is to specify a custom Wasm interface for every type and function. This is more labour-intensive, as it requires developers to write glue code for every language and function by hand.

One project that builds upon this approach is the WebAssembly Component Model presented in [Section 2.2.5](#). It provides a new language for specifying interfaces and can generate glue code for supported programming languages such as Rust, C, C++, Python, JavaScript, etc. Even though it may look similar to the first

2.2 WebAssembly

solution of using standardized serialization, this approach is specifically tailored to Wasm. Thus it can achieve better performance and efficiency.

Another related challenge are common features of the operating system used by almost all programming languages such as file systems, networking or random number generation. Most standard libraries rely on these functions to provide abstractions, e.g. over reading data from a file. An interface for these common APIs must be defined, that is language-agnostic, independent of compilers and compatible across all host platforms, on which the Wasm runtime will eventually execute the Wasm code. A solution generally accepted by compilers and standard libraries is the WebAssembly System Interface (WASI) presented in [Section 2.2.6](#). It also relies on the WebAssembly Component Model for language-agnostic interface definitions.

2.2.5 WebAssembly Component Model

One of the main problems with Wasm is accomplishing a common interface between a Wasm module and the host or between multiple Wasm modules, each with their own specific memory layouts and data representations (see [Section 2.2.4](#)).

The WebAssembly Component Model solves this issue by defining wrappers around Wasm modules, so called *WebAssembly components*[\[12, sec. 2\]](#). On the high-level each Wasm component contains at most one Wasm module and a definition of the module's interface in a new language-agnostic interface language[\[13\]](#). This allows Wasm components to be as modular as regular Wasm modules, but also allows for more interoperability between Wasm components and the host or between multiple Wasm components[\[12, sec. 2\]](#). The common language used to define a Wasm component's interface is called the Wasm Interface Type (WIT) language. Unlike Wasm, it provides various types such as enums, records, option, bit flags, etc. and it allows interface definitions to define new complex types[\[12, sec. 6\]](#).

[Listing 2](#) shows an example for an interface definition in the WIT language. There the world `example-library`, which can be thought of as the interface or world, in which a Wasm component lives in, is defined. It requires a component to export four functions: Three functions for how to add together two signed integers, one function to trim an input string and a function for calculating the sum of a list of unsigned numbers. The fourth function is not exported directly, but instead through an interface, which is able

2.2 WebAssembly

```
package example:example;

world example-library {
  export add: func(a: s32, b: s32) -> s32;
  export trim: func(input: string) -> string;
  export calculate-sum: func(numbers: list<u32>) -> u32;

  export person-greeter;
}

interface person-greeter {
  record person {
    name: string,
    age: u32,
    hobbies: list<string>,
  }

  greet: func(person: person) -> string;
}
```

Listing 2: A WIT definition specifying one world `example-library`, that is used to specify the interface of a Wasm component. The world defines that the Wasm component must export four functions: Three simple functions `add`, `trim` and `calculate-sum` here chosen as examples to display available types and one function `greet`, which is exported through the interface `person-greeter`. The interface combines both a new type `Person` and the `greet` function using it.

to group together type definitions and functions. Here the interface `person-greeter`, which defines the `person` type and a `greet` function is exported by the main world.

The WIT language is designed to be used to specify interfaces for all popular languages such as C, C++, Rust, Python, etc. it defines high-level types, interfaces and worlds, which can be converted from and to low-level Wasm types. The definition for how a specific type of the WIT language is laid out as basic Wasm types is specified by the Canonical ABI also defined as part of the Wasm component model[\[12, sec. 13\]](#).

While the specification of such a canonical ABI is enough for seamless interoperability between Wasm modules created with different technologies, in practice a developer would still have to write glue code for converting between high level types

2.2 WebAssembly

such as enums to low-level canonical ABI types, where an enum would be represented as an integer. In order to reduce bugs and simplify this process of having to write glue code, also called bindings, by hand, projects such as `wit-bindgen` were developed alongside the Wasm component model. They can be used to generate bindings for a specific WIT interface definition to automate the translation between high-level WIT types and low-level Wasm types[\[12, sec. 8\]](#). At the time of writing, `wit-bindgen` can be used to generate bindings for C, Rust, C#, MoonBit, TeaVM-based Java modules and TinyGo-based modules. Generators for other programming languages such as Python and JavaScript also exist: `componentize-py` and `componentize-js`. On the other side, host applications may want to embed Wasm components into their application, where bindings can also be automatically generated (e.g. for the official Wasm runtime `Wasmtime`).

In addition to being able to define language-agnostic interfaces, the Wasm component model provides another feature. It allows multiple Wasm components to be merged into a single new component. In practice this means, that Wasm modules, possibly created with different technologies and languages, each with their own WIT definitions can now be combined into a new component for easier distribution. An example might be a Wasm component that is made up of two smaller components: A component compiled from C code containing highly efficient algorithm implementations and a component containing a Python command-line interface application that depends on the algorithms of the first component.

Precompiled library components such as the one providing implementations for common algorithms in this example, can also be distributed through a dedicated global registry[\[12, sec. 9.4\]](#).

2.2.6 WebAssembly System Interface (WASI)

The WebAssembly System Interface (WASI) is an interface definition between a Wasm module and the host system developed by a subgroup of the WebAssembly Community Group[\[14\]](#). Typically programs communicate directly to the operating system for access to the file system, networking, random number generation, etc. However because Wasm modules are completely isolated from the host system by default, Wasm code has no standardized interface to the Wasm runtime for accessing these parts of the host system. Because Wasm was originally designed for the web, where programs do not have access

2.2 WebAssembly




Figure 3: A plugin system inside of a host application providing the host application methods to access individual plugins. The plugin interface lies between each plugin and the plugin system / host application.

to these features of the underlying operating system, there was no need for it until non-web applications emerged.

WASI solves this problem by defining a common interface similar but not the same as the POSIX standard. It defines its interface based on the Wasm component model in form of separate WIT definitions for every feature set. This then allows programming languages and their standard libraries and compilers to adapt and use these interfaces to support features such as file systems and networking when compiling to Wasm. Note that Wasm runtimes also need to support the WASI specification by exposing the respective functions to Wasm components/modules.

2.3 Plugin systems

In software engineering architectures and patterns define structures used to solve common problems. For example an application can be implemented as a monolithic application or as a layered application that benefits from abstractions. One of these software engineering patterns is the plugin system pattern. It defines plugins as software components, that can implement new isolated features for extension of some host application without modification of the host application itself. On the side of the host application, a plugin system is implemented as a software component for management of and communication to individual plugins.

Figure 3 shows a diagram of all relevant software components for an application with an integrated plugin system. The host application itself runs on the host system

2.3 Plugin systems

and implements a core application. It also contains the plugin system, which is may be hand-written or provided by a library. This plugin system is used by the host application to embed plugins and interact with them. The plugin interface is defined as the common interface, that plugins expect from and provide to the plugin system.

The plugin system architecture is commonly viewed as a software engineering pattern for extending an existing host application with new features. However some research exists, suggesting plugin systems as a general software architecture pattern to build entire applications around[\[15\]](#). This shows that plugin systems are not defined as a single pattern, e.g. compared to patterns such as the factory pattern, but are rather highly variable systems architectures. For example, some plugin systems might allow plugins to depend on each other, while others keep them strictly isolated.

Because of the high variability of plugin systems, the pros and cons also vary. However the main advantages of plugin systems seem to be allowing the extension of host applications with new functionality and being able to enable/disable plugins, depending on whether they are currently required to save system resources. The major drawback of plugin systems is the additional complexity introduces, especially when plugin systems are implemented inside host applications, which where not designed with plugin systems in mind from the start. Incorrect implementations could result in tight coupling between the core host application and the plugin system or unexpected behavior of the host application due to interference with modifications by plugins.

3 Technology comparison of existing plugin systems

To evaluate whether Wasm is a viable technology for versatile plugin systems, one must first understand what criteria make a plugin system good and versatile. This section will perform a technology comparison between several technologies and existing software projects. First, a set of criteria is defined. Then, appropriate technologies and software projects are selected to represent a spectrum of different plugin systems. Next, the technologies and projects are evaluated against the previously defined criteria. Finally, a technology comparison matrix is used to summarize and visualize the results.

3.1 Definition of criteria

In this section criteria for good plugin systems are defined. Each criterion will define a scale from 0 to 5, along with requirements for each score. This scale will be useful later to enable an objective evaluation and comparison of technologies and projects.

3.1.1 Performance

A computer's performance usually refers to the speed it is able to execute software at. For interactive computer systems one generally wants every piece of code to run as fast as possible to minimize its time on the CPU.

In the context of plugin systems, performance also refers to the speed at which software is executed. The three relevant software components necessary to define performance of a plugin system technology are the host system, the plugin system and the plugins managed and called by the plugin system. For this work we define performance as the property that describes how quickly a host system can temporarily transfer execution to a plugin system, which then loads and invokes a plugin's function.

While performance can be measured quantitatively through benchmarks, in practice this is quite hard for plugin systems. Plugin systems and their technologies often vary between host applications as they are by nature highly individual. To benchmark different plugin systems one would have to implement a variety of algorithms and scenarios for a variety of plugin systems. Then one could measure the time each plugin system and plugin takes to execute.

Due to time-constraints and the broad spectrum of knowledge about programming languages and host applications necessary, this work does not use quantitative benchmarks to measure performance. Instead performance is judged through educated

3.1 Definition of criteria

guesses based on benchmarks and comparisons already available for the technologies chosen and built upon by plugin systems.

- 0 – Very slow** The transfer of execution to the plugin systems and invocation of a plugin is highly inefficient. Thus the plugin system is not viable for use within interactive software. Reasons for significant bottlenecks might include heavy serialization or expensive VM-based sandboxing.
- 1, – Slow** Both the plugin system and plugins run very slowly. Transferring execution between the host system and a plugin is inefficient. Therefore this plugin system is also not recommended for use in interactive software, unless these inefficiencies can be somewhat mitigated, e.g. by offloading to other threads.
- 2, 3 – Acceptable** The plugin system and plugins are not fast but their performance is acceptable for interactive software such as text editors. They can negatively impact the user experience by causing stuttering or slow loading times, but there are workarounds to minimize the impact of these problems.
- 4 – Fast** Transferring execution to the plugin system and/or executing a plugin is fast, with only a small overhead, not noticeable by a user. While there is still a small overhead present, it is usually negligible in practice, except in scenarios with real-time requirements.
- 5 – Optimal** Transferring execution and invoking a plugin is virtually instantaneous. There is no measurable overhead. All plugin code executes as fast as if it were implemented natively within the host system.

The scoring outline presented here is intentionally not very specific, without any hard lines between the different scores. It is meant to give only a rough guideline for evaluation, which then needs to be done very carefully on a case-by-case basis. For example, one could evaluate plugin systems based on whether plugins are compiled/interpreted or how large and thus slow plugins might be to load.

3.1.2 Plugin size

The plugin size property refers to the average size of a plugin for a specific plugin system technology. This property does not refer to the size of one specific plugin, but rather it is used to compare different plugin system technologies and how compact and small plugins for them are generally. The average plugin size may vary from technology

3.1 Definition of criteria

to technology due to factors such as static vs. dynamic linking of libraries or the size of the specific language's standard library.

The importance of plugin size depends on the specific use case and user requirements. For text editors specifically a smaller plugin size might result in faster startup times and less time spent downloading or updating the plugin. However reducing the overall impact a program has on the system's resource usage is generally preferred.

Also note that this section only refers to the plugin size and not the size of the entire plugin system. While the plugin system's size is also very important for the memory impact of the host application, it is harder to measure. This is due to the fact that plugin systems are usually very tightly coupled with the host application. To measure a plugin system's size, one would have to disable the plugin system of some host application, without breaking the host application itself. Only then would it be possible to compare system resource usages between the host application with and without a plugin system. Due to the high complexity, the plugin system's size will not be taken into account in this work.

The following scores will be used to evaluate a plugin's size. They are chosen specifically for plugin systems for text editors.

5 – Minimal(< 5KB) Plugin sizes are as minimal as they can possibly be. Plugins contain the minimal amount of program code necessary to achieve their desired functionality. The plugin code format is also made to be very space-efficient, which could be implemented for example through compression and hacks on the byte/bit levels.

4 – Negligible(< 500KB) Plugin sizes are so small that they are negligible in practice. There is no replication of similar information between multiple plugins such as statically-linked libraries.

3,2 – Moderate(< 50MB) Plugins are not very small, however their size is still quite manageable in the context of text editors. Examples could be plugin system technologies, that require a large fully-self contained runtime to be shipped with every plugin. Plugins might also have to contain all libraries and standard libraries that they depend on.

1 – Large(\leq 500MB) Plugins are unusually large specifically in the context of text editors. They are not as easy to manage and during runtime they might also have a non-negligible impact on the memory footprint of their host application.

3.1 Definition of criteria

0 – Very large(> 500MB) Plugins are very large. This may be due to the fact that their internal program logic requires very costly operations such as the virtualization of an entire environment, that must be completely self-contained in the plugin code.

3.1.3 Plugin isolation

Often times plugins contain foreign code. This is especially true for text editors, where plugins are often downloaded from a central registry, also known as plugin/extension marketplaces. This means that plugins downloaded from such sources are usually not validated and thus should be treated as foreign code. Even though there might be checks in place for malicious contents, foreign code should not be trusted to not access its host environment unless otherwise allowed.

For this work we define the property of plugin isolation to describe how isolated a plugin's environment is from its host environment. While there has to be some kind of interface between both environments to make plugins accessible and usable, this interface must not be considered when evaluating plugin isolation. Instead we define plugin isolation completely separated from the interface, meaning if an interface is unsafe by nature, plugin isolation is not automatically violated.

0 – No isolation, required elevated privileges Plugins are not only not isolated from the host application, they also require certain elevated privileges, which the host application usually does not require by itself.

Worst case: Elevated privilege access to current system.

1 – No isolation Plugins are not isolated from the host application. They inherit the host application's privileges without any attempt of the host plugin system to restrict these permissions.

Worst case: Plugins gain the same privileges as the host system, usually this means access to the current user's system and peripherals.

2, 3 – Restricted isolation Plugins are not isolated from the host application by design. Normally the plugin would inherit the host application's privileges, however the host application makes an attempt to restrict plugins from accessing certain critical functionalities. Some examples for restrictions on Linux systems could be allowing only a specific subset of syscalls through `seccomp(2)` [16] or using namespaces (7) [16] to isolate and limit resources. However both of these examples do not use a sandboxing strategy for isolation.

3.1 Definition of criteria

Because these restrictions can come in a various shapes and forms each based on different technologies, during evaluation either 2 or 3 can be chosen as a score. *Worst case: Plugins have similar privileges to that of the host application, except for those specifically disallowed. However this restriction might be able to be circumvented.*

4 – Fully sandboxed, dynamic interface Plugins generally run completely isolated. However their interface is not statically defined, which can lead to vulnerabilities of the interface during runtime due to higher complexity and risk of bugs. Imagine a scenario where a host application exposes only a single interface for passing serialized messages back and forth with a plugin. Then the host application has to serialize and deserialize those messages during runtime. For complex systems, where advanced concepts such as additional shared memory between the host and plugins are used, this interface can become susceptible to logic bugs due to the dynamic interface.

Worst case: Access to parts of the host application not meant to be exposed due to a bug in the interface.

5 – Fully sandboxed, static interface The plugin runs fully sandboxed. It has no way of interacting with the host system, except for statically checked interfaces. Here statically checked interfaces refers to interfaces, that can be proven safe during compilation (or alternatively development) of the plugin system. One way to achieve this might be an interface definition in a common interface definition language. This restriction was chosen because it disallows plugin systems giving full access to parts of a host application without a proper interface definition.

Worst case: Indeterminable, a major bug in the sandboxing mechanism is required.

3.1.4 Plugin portability

Portability stems from the field of distributed systems. A high portability refers to software components, that can be moved from one distributed system A to another distributed system B without having to make any modifications[17]. This assumes that both systems A and B share a common interface for interaction to the software component[17]. In the context of plugin systems for text editors, portability can be interpreted in one of two different ways:

3.1 Definition of criteria

1. Every individual plugin is seen as a software component. This plugin is portable, if it can be loaded into two instances of the same text editor running on different platforms.
2. The entire plugin system itself is seen as a modular software component of a text editor. It is portable if it can be integrated into different text editors and run across different platforms.

This work considers only the first scenario, in which portability refers to each individual plugin, because this scenario is less extensive and the portability of individual plugins is easier to measure. The following scores are used to measure plugin portability:

0 – Not portable The plugin is not portable between different platforms. It is theoretically and practically impossible to run the plugin on different platforms.

1 – Theoretically portable The plugin is theoretically portable between different platforms. In practice this might be very complex and costly, e.g. having to run each plugin in its own dedicated virtual machine.

2,3,4 – Portable with a runtime The plugin is portable between different platforms, but it requires a runtime on the target platform. Because these runtimes can vary from one plugin system to another, a score range from 2 to 4 is specified here. During evaluation the specific runtime has to be analyzed regarding its complexity and impact on the host system. A more lightweight runtime could also enable higher portability of the plugin system itself as described in [Section 3.1.4](#).

5 – Portable by design The plugin is portable between different platforms without requiring a runtime on each host application. In practice this is very hard to achieve. Advanced technologies such as fat binaries, which are binaries that encapsulate compiled machine code for multiple different architectures, might be necessary.

Note that the most extreme scores 0 and 5 are very unlikely for any imaginable plugin system. 0 requires a plugin not to be portable at all, while 5 requires that a plugin is portable to different platforms and architectures which is near impossible to implement on a technical level.

3.1.5 Plugin language interoperability

Text editors and their plugin systems are highly individual software. Some users have personal preferences e.g. keybindings, macros or color schemes, while others may

3.1 Definition of criteria

require tools such as language servers for semantic highlighting and navigation, or tools to compile and flash a piece of software onto an embedded device.

A lot of times plugin systems are used to overcome this challenge of high configurability. Major text editors and IDEs such as Neovim, VSCode, Emacs, IntelliJ or Eclipse provide plugin systems, some even in-built plugins. The advantage of implementing features as plugins is the reduced code complexity and size of the host application. Also providing a plugin systems with a publicly documented interface allows every user and developer to implement their own plugins for their own needs.

This property describes how interoperable plugins written in different languages are. In other words, the larger the set of languages is, in which a plugin can be written for a given plugin system technology, the better its plugin language interoperability is. The more languages are available, the less effort it requires for users to develop their own plugins without the need of learning a new (possibly domain-specific) programming language.

One could argue, that supporting a variety of different languages can result in higher interface complexity and less adaptability to new changes. Even though the complexity and adaptability of interfaces is another important property, which deserves its own rating, it will not be covered in this work.

0 – Domain-specific custom language Only a domain-specific language can be used to write plugins in. This language is specifically designed for given plugin system technology, which is why it is the least interoperable and might be the most unfamiliar and hardest for developers to learn and use.

1 – One language A single general purpose programming language is supported to write plugins in. There might be some developers who are already familiar with the language.

2 – Multiple languages A set of multiple languages is supported to write plugins in. The plugin system is able to abstract over multiple different plugin languages, so that the host application has only a single interface to communicate with plugins regardless of their specific language used.

3 – Build target for some languages The plugin system supports a compilation target for plugins, that is targeted by multiple compilers for multiple programming languages. For example the Java bytecode is a compilation target for the Java, Kotlin and Scala compilers.

3.1 Definition of criteria

4 – Build target for a variety of languages The plugin system supports a compilation target for plugins, that is targeted by a variety of different languages. This score differs from the previous score, that this score’s compilation target is targeted by a considerably higher number of languages with more differences between them. Differences between languages could include dynamic vs. static typing, weak vs. strong typing, interpretation vs. just-in-time compilation vs. ahead-of-time compilation.

5 – Universal build target The plugin system supports a compilation target to which most software can be compiled directly, or which can be embedded indirectly by a compiled runtime. For example all source code is eventually compiled to native ISA instructions specific to some hardware and platform. Thus it is also theoretically possible to package source code such as Python or JavaScript source code and combine it with their specific runtimes inside a native plugin.

3.2 Choice of technologies & projects

An important step during a technology comparison is the correct selection of technologies. When choosing technologies, one has to be careful to not introduce any bias towards certain technologies. This work mainly focuses on plugin systems for text editors, so text editor plugin systems will be the majority of technologies. However it could also be interesting to compare text editor plugin system technologies to plugin system technologies for other applications.

To make an appropriate decision for text editor projects, the popularity of different text editors will be as a metric for selection. As a source for the most popular text editors and integrated development environments (IDEs), the StackOverflow Developer Survey 2024 will be used[\[18\]](#). The platform StackOverflow is known as a forum for developers helping each other by asking questions and exchanging information regarding various technical topics such as programming languages, certain APIs, technologies, etc. Annually it organizes a survey open to all developers regardless of their background to gather statistics about used programming languages, salaries of developers, used development tools, etc. In 2024 a total amount of ~65.000 developers took part in this survey, which is why it can be considered fairly representative. The four most popular text editors and integrated development environments (IDEs) according to the survey are the following. They are described and also considered for evaluation in this work:

3.2 Choice of technologies & projects

1. Visual Studio Code Visual Studio Code (abbreviated as VS Code) is a text editor designed for software development. It provides built-in basic features such as a terminal, version control or themes⁴. However it does not provide built-in integration for specific programming languages or technologies. In VS Code a lot of features are instead packaged as plugins (here: extensions) written in JavaScript which can be installed from a central marketplace. VS Code itself is also written in JavaScript. It is chosen for evaluation in this work, because it has successfully implemented a plugin system able to integrate a broad spectrum of plugins ranging from simple plugins for coloring brackets to highly complex plugins such as programming language integrations or plugins enabling interactive jupyter notebooks.

2. Visual Studio (not evaluated here) Visual Studio is an IDE, that describes itself as “the most comprehensive IDE for .NET and C++ developers on Windows”⁵. It is not to be confused with Visual Studio Code, as Visual Studio provides more built-in features for developing, debugging, testing and collaboration between developers⁶. It also allows users to install plugins (here: extensions) written in C# from a central marketplace. However due to personal unfamiliarity with this IDE and the C# programming language together with the .NET ecosystem, this IDE will not be evaluated in this work.

3. IntelliJ-family IntelliJ IDEA is an editor made by JetBrains for development of JVM-based languages such as Java or Kotlin⁷. According to the StackOverflow developer survey, it is ranked as the third-most popular IDE.

IntelliJ’s so called “Community Edition” is freely available and open-source, however JetBrains also develops proprietary alternatives for other use cases. In fact most of these alternatives such as CLion for C/C++ development or WebStorm for web development are also based on the IntelliJ framework. There are also third-party IDEs built upon the IntelliJ IDEA such as Android Studio for development of Android apps.

IntelliJ provides a plugin system where plugins can be written in Java or Kotlin, that can be used from all IDEs based on the IntelliJ framework. Thus the IntelliJ

⁴<https://code.visualstudio.com/>

⁵<https://visualstudio.microsoft.com/>

⁶<https://visualstudio.microsoft.com/vs/features/>

⁷<https://www.jetbrains.com/idea/>

3.2 Choice of technologies & projects

IDEA is not listed as a single IDE here, but rather a family of various IDEs all based on the same framework using the same plugin system technology.

4. Notepad++ Notepad++ is an open-source text editor focused around minimalism and efficient software⁸. The editor itself is written in C++ with exclusive support for the Windows operating system. It provides a plugin system for loading plugins in the form of dynamically linked libraries (DLLs)[19].

Besides these text editors and IDEs, this work also evaluates plugin system technologies for other application types. Multiple projects and technologies were considered, however due to their similarity to already chosen technologies, missing documentation and time-constraints for this work most of them were disregarded:

VST3 for music production During music production in a digital audio workstation, producers use plugins to simulate a variety of different audio effects, entire instruments or traditional analog devices such as compressors. One notable standard used for these kinds of plugins is the open and free Virtual Studio Technology (VST) 3 standard developed by Steinberg[20]. It specifies a standard for technology on how to implement a plugin system technology between a digital audio workstation (DAW) as the host application and a plugin which applies some effects or produces an audio signal. For example a DAW could send an audio stream to some plugin, then the plugin could apply an effect such as a simple equalizer on the audio signal and return the signal back to the DAW.

On the technical level VST3 plugins are DLLs on Windows, Mach-O Bundles on Mac and packages on Linux. The term package on Linux can have different meanings depending on the Linux distribution and package manager used.

The company behind the VST3 standard also provides a software development kit (SDK) for the C++ programming language and an API for the C programming language.

Microsoft Flight Simulator (not evaluated here) Microsoft Flight Simulator is a simulator for aircraft with focus on ultra-realism⁹. It allows for a wide variety of aircraft, airports and systems such as advanced flight information panels. This is achieved by providing multiple software development kits (SDKs) for plugins (so

⁸<https://notepad-plus-plus.org/>

⁹<https://www.flightsimulator.com/>

3.2 Choice of technologies & projects

called add-ons). These SDKs support plugin languages such as C, C++, .NET languages, JavaScript or WebAssembly[21, sec. [Programming APIs](#)].

However this software project will not be evaluated in this work. It's source code is not publicly available, which makes analysis of its architecture harder. Also a paid license is required to make an appropriate evaluation possible.

Eclipse(not evaluated here) Eclipse is an IDE used during software development for a variety of programming languages. It features a plugin system where plugins are written in Java and a central marketplace for installing plugins. However Eclipse is not chosen for evaluation, because its plugin system is too similar to IntelliJ's. Also IntelliJ is more popular as an IDE [18] and thus required to support a greater variety of plugins.

3.3 Evaluations of technologies & projects

3.3.1 Visual Studio Code

Performance In terms of performance Visual Studio Code performs reasonably well.

It builds on web technologies, specifically Electron which utilizes the Chromium browser engine together with Node.js. VS Code is written in JavaScript, which still imposes some limitations compared to native applications such as pauses due to garbage collection or the single-threaded nature of Node.js modules[22, sec. [Examples.Multithreading](#)]. While JavaScript engines such as the V8 engine try to mitigate most issues, the performance of VS Code still remains slower and more memory-hungry than native alternative IDE and text editor applications written in C, C++, Rust, etc.

Even though the overhead imposed could be considered small and VS Code and it's plugin system is quite fast for today's standards, its performance is still rated with a score of 3. This score is chosen, because VS Code performs below average but still acceptable in comparison to typical native applications.

Plugin size The average plugin size for VS Code plugins is highly variable. Most minimal plugins typically consist of JavaScript or TypeScript code and a JSON manifest file. However, more complex plugins can be significantly larger, for example plugins that bundle third-party libraries and dependencies or language servers.

3.3 Evaluations of technologies & projects




Figure 4: Plugin size distributions of selected VS Code plugins.

Because VS Code is based on web technologies, its plugins can also be used in web environments, where its dependencies have to be bundled, which can further increase plugin sizes. Though on average, the size of VS code plugins a lot due to the size of libraries and resources needed for specific use cases.

Figure 4 shows plugin sizes of various arbitrarily selected VS Code plugins. The largest plugin here is the rust-analyzer plugin with a size of 43MB. The smallest is the Vim emulation plugin with 4MB. The diagram also shows the distribution of different data types and their contribution to each plugin's total size. Here one can see that some plugins contain mostly other data, such as the rust-analyzer plugin which contains the binary rust language server. While other plugins make use of third-party packages such as external dependencies or libraries, such as the SCSS Formatter plugin, there are also plugins with a high share of source code such as the Vim emulation plugin. However for the evaluation of the average plugin size for the VS Code plugin system, other data such as resources, images and binaries must be left out, because this data varies from one plugin to another. Looking at the diagram, one can see, that the size of source code around the high KBs to the low MBs. Thus this plugin system technology is evaluated at a score of 3.

Plugin isolation

VS Code provides a moderate level of plugin isolation. It is based on web technologies such as JavaScript. A lot of web technologies like JavaScript are designed to enable safe code execution on a client's browser. This is achieved through

3.3 Evaluations of technologies & projects

sandboxing, which means running an application in an isolated “sandbox”, where the application cannot escape this sandbox.

VS Code runs all plugins (officially known as extensions) in separate extension host processes to isolate them from the main application[23, sec. [Extension runtime security](#)]. Also plugins do not gain direct access to the main editor’s state such as the DOM tree. Instead only specific APIs are exposed to the extension host.

However the extension host inherits the same permissions as the original VS Code application. This means that plugins running together in a process with the extension host may have full access to filesystems, peripherals or running processes. Because of this inheritance of permissions and the dynamic and flexible nature of JavaScript itself, malicious plugins are not necessarily restricted from accessing the main editor state.

Thus VS Code’s plugin isolation is evaluated at a score of 3. Plugins run isolated from the main application, however malicious plugins may be able to circumvent this isolation through the inherited host privileges.

Plugin portability VS Code plugins are very portable. Since they are build on web technologies, that are required to be portable, plugins can run on various platforms, operating systems and architectures mostly without modifications.

That said, plugins that rely on native binaries such as language servers, which are often compiled in advance and then shipped with the plugin, may require separate build for different platforms, reducing portability.

In summary the portability of VS Code plugins is still very high, except for some edge cases, relying on native code, resulting in a plugin portability score of 4.

Plugin language interoperability Plugin language interoperability describes how many programming languages may be used to develop plugins. VS Code only allows JavaScript or TypeScript, which builds on JavaScript with a type system, as plugin languages.

- asm.js is subset of JS: acts as a compilation target for C/C++ with Emscripten
- Also plugins may include native code. This allows them to ship compiled programs written in C, C++, Rust or entire compiled runtimes for languages such as Python.

3.3 Evaluations of technologies & projects

In such scenarios a plugin's JavaScript code could contain only glue code to forward function calls between the plugin core written in C, C++, Python, etc. and the VS Code host application.

While it is possible to write parts of VS Code extensions in other languages than JavaScript or TypeScript, the complexity of such plugins may increase rapidly. However following a strict evaluation for VS Code's plugin system, it receives a score of 1, because the plugin system always requires plugins to contain some JavaScript code, even if it is just glue code.

3.3.2 IntelliJ-family

Performance IntelliJ IDEs also perform reasonably well. Both the IDE, as well as its plugin system and plugins are written in languages targeting the Java Virtual Machine (JVM). That is, languages such as Java, Kotlin or Scala, which are optimized and compiled to Java bytecode by their respective compilers. This Java bytecode can then be executed by the JVM. During execution the JVM can apply additional optimizations and generate native machine code for execution for example through optimizing Just-in-time (JIT) compilation, only then when it is needed.

Even though JVMs are highly complex systems with loads of mechanisms for optimization, they still introduce some overhead. Also they are responsible for holding the state of programs during runtime and for garbage collecting no longer used object instances.

Thus the performance of plugin systems and plugins based on the JVM, will not be able to outperform native non garbage-collected approaches. This plugin system technology, as it is used in IntelliJ-based IDEs is rated with an average performance score of 3.

Plugin size Plugin sizes for IntelliJ plugins also vary a lot. IntelliJ plugins are Java Archive (JAR) files, that can contain source code in the form of class files, resources or any arbitrary data. Small plugin JARs might contain only a few source files, while larger plugins might rely on various dependencies, which are all compiled and statically linked into the plugin JAR file.

For IntelliJ plugins an extensive analysis of plugin size distributions was not possible due to the complexity of internal plugin structures. However some plugins

3.3 Evaluations of technologies & projects

were selected as examples for their total size: The official plugin for Rust language support has a size of 54MB, the official Python language support plugin 69MB and the Prettier formatter plugin has 508KB. Thus the average plugin size for the IntelliJ plugin system is evaluated at a score of 3.

Plugin isolation Both the IntelliJ IDEs and plugins are written in JVM-based languages. No documentation or information was found on the security of IntelliJ plugins. Thus it is assumed, that IntelliJ plugins are loaded into the same JVM as the host IDE, inheriting the host's privileges.

Because there are no isolation mechanisms in place, the plugin isolation of plugins for IntelliJ-based IDEs are evaluated at a score of 1.

Plugin portability Since IntelliJ plugins are written in Java or other JVM-based languages, they benefit from Java's cross-platform portability. This allows plugins to run consistently across different environments. Because IntelliJ itself is build on the same technology, this enables every plugin to run on every platform, as long as plugins do not introduce platform-specific behavior such as relying on native binaries. Plugin portability is rated at a score of 4 for all IntelliJ-based IDEs, as Java bytecode is portable because the JVM is still required as a runtime on the target platform.

Plugin language interoperability IntelliJ supports all plugins that can be compiled to Java bytecode, as long as they adhere to the IntelliJ Platform API, which can vary from one IntelliJ-based IDE to another.

While only Java and Kotlin are officially supported, other JVM-based languages such as Groovy or Scala could also be used, as they also compile to Java bytecode. This plugin system, which supports one compilation target used by a variety of languages, is evaluated at a score of 4.

3.3.3 Notepad++

Performance Notepad++ is written in C++ and compiled to native machine code for the Windows operating system exclusively. It tries to maximize efficiency and minimize the impact on the system it is running on. Its plugin system is based on compiled dynamically linked libraries (DLLs). A plugin developer compiles their plugin, which can be written in any arbitrary language, to a DLL, which is a library containing machine code and lists of imported and exported symbols. Notepad's

3.3 Evaluations of technologies & projects




Figure 5: Plugin size distributions of selected Notepad++ plugins. They were analyzed by running `winedump` on a Linux system on the plugin DLL files. Then the size of code and size of initialized data fields were extracted to show here.

plugin system can then load these libraries at runtime via the `LoadLibrary` Win32-API call and execute arbitrary functions exported from the plugin.

This is the fastest way for how host applications can embed plugins. It relies only on the operating system for loading already compiled machine code at runtime. There is no additional overhead except having to load the DLL itself into memory. Thus Notepad++'s plugin system is evaluated at a score of 5 for its optimal performance.

Plugin size Notepad++ uses DLLs as its plugin format, meaning plugins are compiled directly to machine code. These binaries may include statically linked libraries and dependencies, though some exceptions for example for `libc` might exist. [Figure 5](#) shows selected plugins and their total sizes and size distributions. The size distribution into machine code, resources and other data was extracted from the respective DLLs through use of `winedump` on a Linux system. The bar chart shows, that most data inside each plugin comes from machine code and not from any other resources or binaries. Given this analysis and plugin sizes being under 500KB for all plugins, a rating of 4 is assigned to this plugin system technology. This high rating is due to the compactness of machine code, which generally is generally minimized by compilers for better runtime performance already. In practice smaller binaries require less data transfer rates between the memory and the CPU and thus result in higher cache hit rates and better performance.

Plugin isolation However Notepad++'s plugin isolation is rated poorly at a score of 1. Since Notepad++ loads plugins as dynamic libraries at runtime, these libraries

3.3 Evaluations of technologies & projects

inherit all privileges from the host editor. There is also no sandboxing in place, so plugins essentially run on the exact same level as the host editor. This is a severe risk, because the user has to trust both the plugin developer and also has to verify the integrity of a downloaded plugin.

Plugin portability Plugin portability for Notepad++ plugins is rated at a score of 1. Notepad++ itself supports only Windows 8.1 - 11 as an operating system for multiple architectures¹⁰. Plugins have to be compiled and shipped for every supported architecture separately. Thus plugins are compatible in theory, because one could run plugins compiled for another architecture inside a virtual machine packaged as a plugin itself. However this is not feasible in practice due to performance losses and the complexity involved.

Plugin language interoperability With native machine code as plugins, Notepad++ allows for great plugin language interoperability. In practice it common to write plugins in languages that compile directly to machine code, such as C, C++ or C#. Notepad++ provides official header files for the API to Notepad++ for plugin development in C++. There are also community-driven API definitions and templates for plugin development in Ada, C#, D, Delphi or Ada available[19].

One can also argue, that every piece of software running today is eventually compiled to or interpreted as machine code. Thus it may be possible to embed higher-level technologies such as a Python runtime along with plugin logic written in Python inside a plugin.

Because native machine code is a universal build target for all languages and technologies and the resulting versatility, this plugin system is rated at a score of 5 in terms of plugin language interoperability.

3.3.4 VST3

Performance The VST3 standard for plugins that create or process audio relies on native machine code just like Notepad++. However it's file format also accommodates for the fact that plugins might be run on more than one platform. Internally a VST3 file embeds either a DLL for Windows plugins, a Mach-O bundle for MacOS plugins or generic packages for Linux plugins.

¹⁰https://github.com/notepad-plus-plus/notepad-plus-plus/blob/master/SUPPORTED_SYSTEM.md

3.3 Evaluations of technologies & projects

During runtime a host application has to check whether a VST3 plugin contains machine code compiled for the current architecture. Then it is able to link the machine code at runtime through the operating system just like Notepad++ does.

While there is a small overhead of checking if the targeted platform of a plugin is correct for loading a plugin, there is no overhead during execution of plugin code. Thus the VST3 technology is evaluated at a score of 5.

Plugin size Because VST3 plugins are distributed as platform-specific binaries, their size can also remain relatively small comparable to Notepad++'s plugins. On the other hand, VST3 plugins often contain more resources such as images for providing the user with a graphical user interface.

Depending on the specific use case, one could rate VST3 plugin sizes at a score of 3 or 4. However to enable objective comparison and VST3's strong similarity to Notepad++, plugin size is rated at a score of 4.

Plugin isolation VST3 plugins are loaded as dynamic libraries just into the host's memory space inheriting privileges comparable to Notepad++'s plugin system. There is no isolation between the host application and VST3 plugins, which can be used by malicious plugins posing a significant risk to the end-user. Thus plugin isolation is rated at a score of 1.

Plugin portability VST3 plugins itself are not portable. They are compiled for every architecture and operating system and then distributed separately. Even though the VST3 documentation states that the source code of plugins can be reused across multiple platforms, the plugin portability stays unaffected as this benefits just the plugin developer. The portability of VST3 plugins is therefore rated at a score of 1.

Plugin language interoperability VST3's plugin language interoperability is also similar to that of the Notepad++ plugin system technology. It's plugins are also based on native machine code, and thus it is possible in theory to compile or embed most technology into a plugin. However in practice it is more common for plugin developers to develop plugins with C++ using the official VST3 software development kit. The universality of the native binary format earns VST3 a rating of 5 for plugin language interoperability.

3.3 Evaluations of technologies & projects

	Performance	Plugin size	Plugin isolation	Plugin portability	Plugin language interoperability
Visual Studio Code	3	3	3	4	1
IntelliJ-family	3	3	1	4	3
Notepad++	5	4	1	1	5
VST3	5	4	1	1	5

Table 1: Technology comparison matrix for selected technologies and software projects.

3.4 Results of the technology comparison

Three different plugin systems and one standard for a plugin system technology have been evaluated with the previously defined criteria. The results can be seen in a technology comparison matrix in [Table 1](#). Here every technology row is evaluated across all five criteria, namely performance, plugin size, plugin isolation, plugin portability and plugin language interoperability. Each criterion is rated on a scale from 0 (very poor) to 5 (excellent), however the smallest score that can be found here is 1.

Visual Studio Code's plugin system performs with average scores of 3 across performance, plugin size and plugin isolation. It provides good plugin portability from its usage of web technology, but is not able to provide a lot of plugin language interoperability due to the limitation to JavaScript/TypeScript. The plugin system used in all **IntelliJ-based IDEs** achieves similar scores for performance and plugin size, however it does not provide any plugin isolation. It is able to achieve the same plugin portability as VS Code and a higher score for plugin language interoperability because it allows any JVM-based languages for plugins. **Notepad++** is another text editor with efficiency in mind. This shows, because it achieves the highest score for performance with plugin sizes that are negligible in practice. However it provides terrible plugin isolation and plugin portability due to plugins consisting of natively compiled machine code. On the other hand, native machine code also acts as a universal compilation target, allowing pretty much every other technology to be compiled into a plugin. This allows

3.4 Results of the technology comparison

for an optimal plugin language interoperability. **VST3** is a standard for a plugin system technology. It follows the same approach as Notepad++'s plugin system with native machine code inside its plugins. Thus it achieves the same scores, although in practice the VST3 format is merely a wrapper around native machine code, providing additional features not considered here.

Now a critical discussion of the chosen criteria is presented, outlining possible improvements and other relevant criteria, not covered in this section. Some criteria, namely the average plugin size and performance, have been analyzed only qualitatively due to time constraints. However, these criteria can and should be analyzed quantitatively to provide more accurate and verifiable results. There are also criteria that may also be important for plugin systems in general, not covered here:

Complexity and flexibility of plugin system interface Plugin systems with more flexible and dynamic interfaces without a fixed interface definition may enable faster development and testing of new features.

Developer experience Developer experience could be used as a measure for how easy it is for developers to develop plugins.

Debuggability Some technologies might be easier to debug and inspect during runtime. This allows plugin developers to spend less time searching for bugs, which then results in a faster plugin development.

Plugin system size This section analyzed the average size of plugins. The size of entire plugin systems was not considered. However when plugins get smaller, the size of plugin systems could increase, due to logic being outsourced from plugins to the plugin system. Thus to optimally evaluate sizes, one should consider both the plugin size and plugin system size and not just rely on one.

To conclude, four plugin system technologies were compared and the results visualized in a technology comparison matrix. While the matrix shows the strengths and weaknesses of these technologies across different criteria, there is more to consider when choosing an appropriate plugin system technology. In the future, additional criteria should be defined for a more accurate comparison.

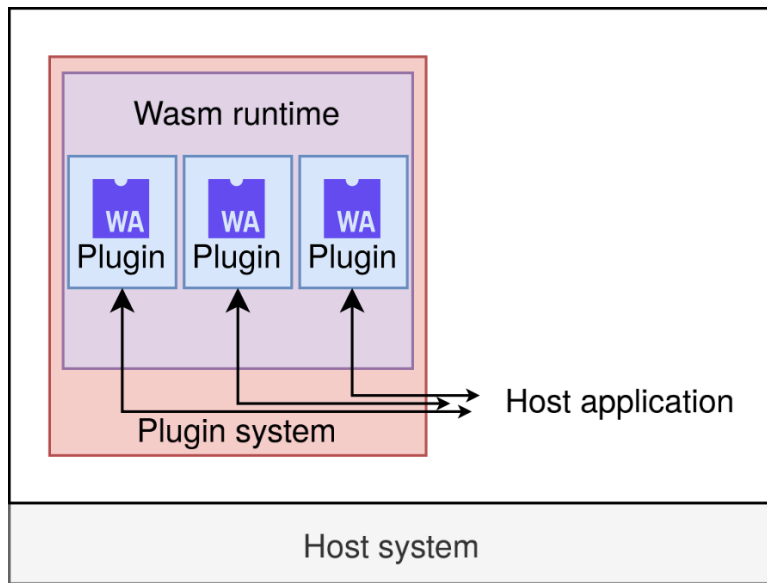


Figure 6: A plugin system using a WebAssembly runtime to embed WebAssembly plugins inside of a host application.

4 WebAssembly for plugin systems

This section explores Wasm as a technology for how to run individual plugins in a plugin system. Wasm is able to provide performance ([Section 2.2.3.A](#)), safety through sandboxed isolation ([Section 2.2.3.B](#)), portability ([Section 2.2.3.C](#)) and other features. To find out, whether Wasm is competitive to the previously evaluated plugin systems, Wasm will be evaluated in this section.

First, an overview over the architecture of Wasm plugin systems and the relevant software components is given. Then Wasm is evaluated as a technology for the criteria defined in [Section 3.1](#) and compared against previous plugin systems. In the end related projects implementing Wasm plugin systems and their technology choices are presented briefly.

Generally plugin systems inside a host application manage plugins and allow interaction between the host application and plugins ([Section 2.3](#)). Wasm plugin system architectures follow this approach, however a Wasm runtime is added as an additional layer between the plugin system and Wasm plugins. [Figure 6](#) shows an exemplary plugin system architecture similar to the one shown in [Figure 3](#). However here the Wasm runtime explicitly owns and manages the Wasm plugins, consisting of their Wasm modules and runtime state such as linear memory, the stack, etc. Also the Wasm

4 WebAssembly for plugin systems

runtime and its plugins exist completely inside the plugin system. They are isolated and sandboxed through software, as opposed to other means of isolation such as running plugins in isolated processes.

4.1 Evaluation of criteria

Performance The execution of WebAssembly bytecode is very fast by design. While there is some overhead for loading Wasm plugins and compiling them, at runtime Wasm is able to achieve speeds comparable to native speed by a single-digit factor (see [Section 2.2.3.A](#)). Also Wasm bytecode is designed to be compact (see [Section 2.2.3.E](#)), which can further increase its performance, as smaller data can generally be loaded and operated on by the CPU.

In real interactive applications systems this overhead is negligible in most cases. Thus Wasm is rated at a score of 4 for the performance criterion.

Plugin size With Wasm being originally designed for the web, the size and compactness of its bytecode are very important. On the web larger files result in longer file transfers and loading times, which can possibly impact the user experience negatively. Thus Wasm was designed with compact bytecode. In comparison to natively compiled machine code, it is larger by a factor of around 2.4x (see [Section 2.2.3.E](#)), which is acceptable for text editor plugin systems.

However, note that Wasm plugins are also able to embed higher-level languages that require a runtime such as Python or JavaScript. In these cases a small runtime is compiled to Wasm and is distributed along with the source code of the program written in Python/JavaScript. This may introduce a comparatively large overhead, as multiple Wasm plugin may contain redundant information, e.g. the same Python runtime.

Because Wasm is generally able to achieve relatively compact plugin sizes with only larger plugin sizes for some languages which cannot be compiled to Wasm directly, it is rated at a score of 3.

Plugin isolation The execution of Wasm code is sandboxed by default. It's execution semantics are defined formally by the specification in such a way, that there are no loopholes for Wasm code to escape this sandbox in theory. Access to the host is allowed only through functions, that are explicitly exposed by the hos. By default Wasm runtimes do not expose any of these interfacing functions to Wasm

4.1 Evaluation of criteria

modules. The design goal for Wasm regarding safety is described more detailed in [Section 2.2.3.B](#).

Also Wasm's type system does not allow for any methods to share data between the host and Wasm code other than its simple type system or by having the host access the linear memory of a Wasm module. This means that every interaction between the host application and Wasm plugins must be defined statically. No dynamic sharing of data such as memory mapping or passing pointers is allowed. The only exception for this are funcrefs and externrefs, which are references opaque to Wasm, meaning Wasm code cannot forge and read the bits of these references[8, sec. 2.3.3].

Because Wasm allows for perfectly isolated plugin execution in theory and its type system does not allow dynamic sharing of data, the plugin isolation criterion is rated at a score of 5.

Plugin portability Wasm is designed with portability as a goal. Originally it was designed for the web, where portability is necessary to support various browsers, platforms and architectures (see [Section 2.2.3.C](#)).

For that the execution of Wasm requires a small runtime, able to execute the relatively simple Wasm code. Once Wasm code is compiled, it is then portable to any platform, for which any Wasm runtime exists.

Because of the need of a small Wasm runtime and the extreme portability of Wasm plugins themselves, the plugin portability for Wasm plugins is rated at a score of 4.

Plugin language interoperability Wasm is designed to be a universal compilation target for various languages. One design goal of Wasm is the independence of languages as described in [Section 2.2.3.D](#), so that various languages can be supported equally. The the time of writing, the official Wasm website lists 25 languages and technologies, which can be compiled or embedded in Wasm modules. Examples for popular languages that can be compiled directly to Wasm are Rust, C, C++, Java. Other languages such as Python are not compiled directly to Wasm. Instead the source code written in these higher-level languages is distributed together with minimal runtimes. For example when Python is embedded in Wasm, a Python interpreter is compiled to Wasm instead and distributed along with the original Python source code. This is what the project Pyodide does (<https://pyodide.org>).

4.1 Evaluation of criteria

	Performance	Plugin size	Plugin isolation	Plugin portability	Plugin language interoperability
Visual Studio Code	3	3	3	4	1
IntelliJ-family	3	3	1	4	3
Notepad++	5	4	1	1	5
VST3	5	4	1	1	5
Wasm	4	3	5	4	5

Table 2: Final technology comparison matrix for selected technologies and WebAssembly.

Because Wasm is able to act both as a compilation target comparable to existing compilation targets such as x86-64 or ARM machine code and as an environment for embedding of higher-level languages, it is rated at a score of 5. It allows plugins to be written in various different languages without technology restrictions for plugin developers.

4.2 Evaluation of WebAssembly for plugin systems against previous technologies

The last section [Section 4.1](#) evaluated all five criteria for Wasm as a technology for plugin systems. Now that existing technologies and Wasm have been evaluated, they can be compared. [Table 2](#) shows the previous technology comparison matrix [Table 1](#) with Wasm and its scores as a new row. It is visible, that Wasm performs great overall. On the one hand it lacks performance and has slightly larger plugins in comparison to Notepad++ and VST3. However both of these technologies use native machine code for execution, which achieves the optimal performance, because they do not require a runtime and thus no overhead is present. While Wasm cannot take the first place for these criteria, it is able to outperform previous technologies in plugin isolation and portability while still acting as a universal compilation target. In comparison to the

4.2 Evaluation of WebAssembly for plugin systems against previous technologies

plugin systems used by VS Code and all IntelliJ-based IDEs, Wasm provides severe advantages across all five criteria evaluated in this work.

What stands out the most are the exceptionally low scores for existing technologies for the plugin isolation criterion. Most plugin systems do not provide any isolation, except for the plugin system of VS Code, which takes advantage of JavaScript’s sandboxing capabilities to a certain degree. Wasm however achieves a perfect score of 5 for plugin isolation with its execution sandboxed by default while only allowing explicitly exposed interfaces by the host.

Overall Wasm performs very well compared to existing plugin system technologies. This leads to the question, if Wasm is viable for plugin systems in practice.

4.3 Related projects using WebAssembly plugin systems

This section presents noteworthy software projects, which have implemented a Wasm plugin system. While Wasm itself acts as the fundamental technology to execute individual plugins similar to virtual machines, in practice other choices regarding challenges such as interfacing must be made. These challenges and issues might be solved differently by various plugin system implementations.

Zed, a self-proclaimed “next-generation” text editor, uses a Wasm plugin system. Plugins (officially called extensions) can add features in the form of languages, themes, icon themes, slash commands or context servers. Zed plugins come in the form of Wasm components, taking advantage of the Wasm component model. The main Zed text editor projects defines the plugin interface in the form of a WIT definition, as explained in [Section 2.2.5](#). Currently Zed also provides a small shim library for the Rust programming language. This shim library contains a Rust trait (similar to interfaces in other programming languages) `zed_extension_api::Extension`, which can be implemented to abstract away the WIT interface. With this abstraction, plugin developers do not have to deal with Wasm-specific intricacies. However if developers want to make use of Wasm’s language interoperability, they must generate bindings adhering to the WIT definition themselves, as Zed only supports Rust as the only official plugin language[\[24\]](#).

ZelliJ is an application similar to a terminal multiplexer. It allows to split up one terminal instance into multiple panes or manage terminal sessions. Similar commonly known terminal multiplexers are Tmux, xterm or the Windows Terminal. ZelliJ features

4.3 Related projects using WebAssembly plugin systems

a Wasm plugin system, to allow accessible development of new features in various programming languages[25]. The plugin system is also used by the ZelliJ project itself. Some features such as a status bar, the about page or even the plugin manager itself are developed as plugins, which can be advantageous for developer accessibility and learning purposes. For interfacing between the host application and Wasm plugins, ZelliJ uses Protobuf (<https://protobuf.dev/>). Protobuf specifies a language-agnostic interface definition language not particularly for Wasm systems, as described in Section 2.2.4. An interface definition can then be used to generate bindings for both the host application and Wasm plugins. It also uses the WASI for abstraction of common system interfaces such as the filesystem or networking (see Section 2.2.6)[25].

Extism is a generic library used to implement Wasm plugin systems in applications written in various languages. It is originally written in Rust, but provides bindings in the form of so-called Host Software Development Kits (SDKs) for other host languages. For plugins it provides so-called Plugin Development Kits (PDKs), which provide bindings to the common interface. It uses a custom serialization format to pass data between hosts and Wasm plugins. However users can build on top of this format by using other serialization libraries such as Protobuf or JSON[26].

4.4 Results

Overall, WebAssembly performs at least as good as the four analyzed technologies in 3 of the 5 criteria. It can achieve **perfect plugin isolation** through its sandboxing without complex dynamic interfaces at runtime, while some other plugin system technologies have no mechanism for isolation in place at all (IntelliJ-family, Notepad++, VST3) or provide just partial isolation (VS Code). Wasm provides **very good plugin portability** comparable to the JavaScript-based plugin system of VS Code or the JVM-based plugin system of the IntelliJ-based IDEs. This degree of portability is by design, as Wasm code must run in browsers on all architectures and platforms in a web context. However Wasm's portability is not rated at a perfect score here, because host applications still have to provide a small Wasm runtime for execution. Also Wasm achieves a **perfect plugin language interoperability**. In theory it acts as a universal compilation target for various languages today, however language support is still rapidly evolving and improving over time. Wasm allows various languages such as Rust, C, C++, etc. to be compiled directly to Wasm bytecode or programs written in higher-level interpreted

4.4 Results

languages such as Python to be embedded by distribution of an additional runtime for those languages. In comparison, some other plugin system technologies evaluated in this technology comparison, namely VS Code and IntelliJ-based IDEs, limit plugins to a single language or a compilation target used only by a few languages such as JVM bytecode. Other plugin systems such as Notepad++ or VST3 use native machine code, which acts as a universal compilation target, just as Wasm bytecode.

On the other hand, Wasm performs worse than other technologies for the performance and plugin size criteria. It achieves a **very good, but not optimal performance** at a score of 4. The only two plugin system technologies are those of Notepad++ and VST3, which both use native code as their execution format. While only native code can achieve perfect performance scores, Wasm's peak slowdowns in comparison to native code seem to be around single-digit factors, which still outperforms other plugin systems, namely VS Code and IntelliJ-based IDEs. The scoring for Wasm's plugin size is similar to its performance scoring. Wasm achieves average plugin sizes, slightly larger than native machine code, again varying only by single-digit factors.

Also some existing Wasm plugin system projects were presented briefly. All projects use Wasm as their fundamental technology for code execution, however their technology choices for optimizations and interface are highly variable. While some plugin systems take advantage of the evolving ecosystem around the Wasm Component Model and WASI specification, others rely on third-party serialization libraries such as Protobuf or even specify a completely custom Wasm plugin interface for users to build on top of.

5 Proof of concept: Implementing a WebAssembly plugin system for a text editor

The first part of this work did a technology comparison to find out if Wasm is feasible as a technology for plugin systems. In this second part, a basic proof of concept for a Wasm plugin system will be implemented for a text editor. To make this process as realistic as possible, to learn as much as possible from this implementation, a text editor that is already being used today by developers will be used.

The text editor chosen is the Helix editor (<https://helix-editor.com/>). It is a terminal-based text editor written in Rust with controls similar to Vim, however it provides more features out of the box. Currently the Helix editor is missing an official plugin system, although there are some community-driven projects trying to implement plugin systems.

5.1 Requirements

This section presents the requirements for the plugin system and plugins. It divides the requirements into functional and non-functional, to distinguish between the what the system should do and how it should do it. Also a list of non-requirements is presented to put constraints on the complexity of this proof of concept.

5.1.1 Functional

Plugin metadata Plugins shall contain metadata in some form. This is required, so that users can view information about plugins such as their name, version, description, etc.

Plugin loading Plugins are loaded by the plugin system at startup. Additional dynamic loading of plugins at runtime is not required to reduce complexity.

Modification of editor state Plugins shall be able to modify the program state of the core editor. For this proof of concepts, 2-3 different interface features should be exposed to every plugin.

Invocation through event hooks The plugin system shall allow plugins to register event hooks. Event hooks are functions of a plugin, that get invoked by the plugin system on certain events in the core editor. For this proof of concept 1-2 different event hooks suffice.

5.1 Requirements

5.1.2 Non-functional

Plugins as self-contained WASM files Plugins shall exist only as fully self-contained WASM files. They do not come with additional resources or metadata and require no dependencies during runtime. This limits the complexity of the plugin system, as no inter-plugin dependencies are necessary.

Sandboxed plugin execution The execution of plugins must be fully sandboxed. Plugins may only access those interfaces provided to it by the plugin system. This is a key requirements for a possible implementation of a permission system, where plugins can only access permissions explicitly granted by the user.

5.1.3 Non-requirements

No complex interfaces This proof of concept is not required to allow plugins access to complex interfaces such as GUIs due to time constraints.

No system calls Usually plugins need access to the system such as file system access, networking, or random number generation. These interfaces are defined by the WASI specification. However this proof of concept is not required to provide WASI access to plugins for simplicity.

No production-readiness The implemented system is not required to be production-ready. Thus no time-intensive tasks such as unit testing, coverage, or benchmarking for optimal solutions are required.

5.2 System Architecture

This section describes the architecture and technologies used for the Wasm plugin system and plugins. Even though a minimal proof of concept is implemented, it is still important to choose the technologies used carefully, so that the minimal system can represent more complex Wasm plugin systems as well as possible. This proof of concept is used to determine, whether Wasm and the technologies in its current ecosystem are ready for use and if they can be scaled up to larger software projects.

As the base technology the Wasm Component Model (see [Section 2.2.5](#)) is chosen. It offers great plugin language interoperability, while also building on top of Wasm's rather basic type system. It allows the plugin interface to be specified in the WIT language, which scales well for larger real-world systems, where the plugin interface might become very complex.

5.2 System Architecture




Figure 7: Dataflow diagram for the process of loading *.wasm files into a Wasm runtime and storing it in the main editor state. The instantiated plugins inside the Wasm runtime contain additional data such as the linear memory, tables, or the program stack.

The core editor project then defines a single WIT definition containing the entire interface for plugins consisting of type definitions, imports and exports. Plugins consist of individual Wasm components, which must implement this WIT definition.

Figure 7 shows the dataflow for how plugins are loaded into the editor. A plugin loader component reads the Wasm Component files and loads them into a Wasm runtime. The Wasm runtime contains objects such as the linear memory, tables or a stack for every Wasm module/component. While instantiating plugins, the plugin loader also makes sure that each plugin implements an interface according to the WIT definition. Finally, the Wasm runtime is stored in the core editor state, which is used by helix for providing access to the editor context from most points in the program.

The design for allowing plugins to modify the core editor state and event hooks are rather straightforward. Modification of the core editor state works by exposing functions to the Wasm components by defining them as imports in the WIT definition. Event hooks on the other hand work by having each plugin export one function per event type by defining them as exports in the WIT definition. These functions are then called by the core editor, each time a certain point in the program is reached.

5.2 System Architecture

```
package helix:plugin;

interface types {
  record plugin-metadata {
    name: string,
    version: string,
    description: string,
    keywords: list<string>,
  }
  enum log-level {
    info,
    warn,
    error,
  }
}

world base {
  use types.{log-level, plugin-metadata};

  export get-metadata: func() -> plugin-metadata;
  export initialize: func();

  import log: func(level: log-level, msg: string);
  import set-editor-status: func(msg: string);
  import get-text-selection: func() -> option<string>;
}

world keyevents {
  export handle-key-press: func(input: char);
}

world modify-buffers {
  import write: func() -> result<_, string>;
  import close-buffer: func() -> result<_, string>;
}
```

Listing 3: The entire WIT definition for all plugin interfaces. For more complex interfaces it should be divided into multiple files to modularize feature sets.

5.3 Implementation

5.3 Implementation

```
package component:my-example-plugin;

world my-example-plugin {
    include helix:plugin/base;
    include helix:plugin/keyevents;
    include helix:plugin/modify-buffers;
}
```

Listing 4: A WIT definition for a single plugin. It defines a single world `my-example-plugin` as the plugin’s interface and reexports/includes a subset of the worlds defined in the main package `helix:plugin` defined in [Listing 3](#).

This section provides an insight into chosen parts of the implementation of the Wasm plugin system. It outlines the choices made during implementation, the technologies used for implementation and the challenges faced.

A central part of the implementation is the Wasm runtime. It is used as a library for loading and interacting with Wasm components. For this proof of concept the official Wasmtime runtime is chosen. It is the official reference implementation for the Wasm specification, providing better support for newer features such as the Wasm Component Model than other runtimes. Also it supports WASI, which may be useful for future testing on this project. Both the Helix editor and the Wasmtime library are written in the Rust, which is why this proof of concept is also fully implemented in Rust.

[Listing 3](#) shows the main WIT definition for the plugin interface. It defines a WIT package called `helix:plugin`, which contains one WIT interface and three WIT worlds. The WIT interface is used here to define complex types, which are then imported and used in the worlds. Normally this WIT definition would specify a single world, which contains all the imported and exported functions required for every plugin. However in practice every plugin only uses a subset of all interfaces. This is why the interface is split into three different worlds, all providing different feature sets to the plugin. Here a base feature set called `base` is defined. It must be implemented by all plugins. Additionally two worlds `keyevents` and `modify-buffers` for allowing plugins to hook into key events and for allowing access to the state of currently opened buffers are defined.

With this modular approach of feature sets in place, every plugin can then define its own WIT definition containing a single world. [Listing 4](#) shows a world of such a plugin,

5.3 Implementation

which is now able to simply reexport the worlds of the main WIT definition by using the `include` statement. This allows plugin developers to choose only a subset of features required by their plugin. At runtime, the plugin system then iterates through all worlds and checks whether each world's interface is provided by the plugin. After the developer has written their plugin WIT definition, a tool like `wit-bindgen`, `componentize-py` or `componentize-js` can be used to generate bindings and/or project skeletons, which adhere to the correct WIT interface.

Another challenge during implementation is the storage of plugin metadata inside the Wasm component. Wasm modules allow exporting of immutable or mutable globals, which could be used to export a global variable that contains the name of the plugin, its version, description etc. However the component model, or specifically its WIT language does not allow the specification of mandatory exported globals. Thus a workaround is used, where every plugin has to export a function `get-metadata`, as it can be seen in [Listing 3](#), which returns a record of all plugin metadata. This function can then be called by the host editor to get and display information about a specific plugin.

5.4 Verification and validation

This section contains a verification and validation of the proof of concept. First it verifies whether all requirements specified in [Section 5.1](#) are fulfilled. Then it validates to what extent the implemented Wasm plugin system is a good and versatile plugin system according to the criteria defined in [Section 3.1](#).

5.4.1 Verification

In the following, all functional and non-functional requirements are verified:

Plugin metadata The implemented plugin system requires functions to specify a function, that returns plugin metadata, currently consisting of the plugin's name, a version, description and list of keywords.

Plugin loading The plugin system implemented can load multiple plugins during startup of the Helix editor. However the list of plugins is currently hardcoded in the editor's source code.

Modification of editor state With the WIT definition shown in [Listing 3](#), plugins can make five calls to the editor: Log messages, query the currently selected text, set

5.4 Verification and validation

the message in the editors status bar, write the current buffer's contents to disk or close the current buffer. This shows, that plugins can access the core editors state. In the future this list can be expanded by simply defining more interfaces in the WIT definition and implementing them in the core editor.

Invocation through event hooks The plugin system currently allows only a single event hook. A plugin can hook into are key press events, by including the keyevents world of the main WIT definition in [Listing 3](#) in its own definition, as it is being done in [Listing 4](#).

Plugins as self-contained WASM files For the implemented plugin system, plugins are standalone WASM files. They contain program code, metadata and can additionally contain binary data segments as resources.

Sandboxed plugin execution Plugin execution is completely sandboxed, assuming that Wasmtime as the chosen Wasm runtime does not introduce any loopholes allowing sandbox escape attacks. Plugins can only access interfaces exposed to it by the core editor and it is not able to access other editor state.

5.4.2 Validation

Furthermore one can validate, whether the implemented plugin system is useable in practice. This will be done by doing a rough reevaluation of the criteria defined in [Section 3.1](#) in the following:

Performance There were no measurements for the performance of the implemented plugin system due to time-constraints. However Wasmtime compiles the Wasm bytecode at the startup of the editor, which is why there should be only a minimal overhead present when Wasm functions are invoked e.g. due to event hooks.

Plugin size There were plugin sizes measured for two different plugins, both containing only the minimum required code for a plugin to load. The first plugin is written in Rust with bindings generated by wit-bindgen and then compiled on release mode. It contains close to zero program logic, except for the relatively small amount of plugin metadata. The resulting Wasm plugin has a file size of 17.7KB.

The second plugin contains the same program logic as the first one. However it is written in Python, with its project skeleton generated by the tool `componentize-py`. The resulting plugin size is 35.1MB.

5.4 Verification and validation

The reason for the Python plugin being more than 1900 times larger than the Rust plugin is that Python cannot compile directly to Wasm. Instead a Python runtime, which can be compiled to Wasm, has to be included in the Wasm plugin. Note that the Rust plugin is also unexpectedly large, even though it contains almost no program logic. This is because the Wasm plugin must still specify all imports and exports for the interfaces chosen in its WIT definition. Also according to the component model it must provide an allocator to the Wasm runtime to allow passing of complex data types as arguments for function invocations[\[13\]](#).

Plugin isolation Sandboxed code execution is already defined as a non-functional requirement and fulfilled for this plugin system. Thus it can also achieve perfect plugin isolation in practice, assuming no bugs are present.

Plugin portability Both the Rust and Python Wasm plugins were compiled on a Windows system and then loaded into a Helix instance on a Linux (NixOS) system. Due to time-constraints no additional combinations of platforms and operating systems were tested.

Plugin language interoperability Plugins were written with the Rust and Python programming languages to show that both directly compiled languages and interpreted languages can be used as Wasm plugins.

5.4.3 Conclusion

All in all, the implemented Wasm plugin system is able to fulfill all requirements and also performs reasonably well for the criteria defined in [Section 3.1](#). However a more in-depth analysis of the plugin system regarding most of the criteria might be useful in the future. Overall, Wasm looks to perform quite well as a plugin system technology. Its component model enables quick development without time-consuming tasks like writing bindings by hand. Also it could allow for a simple integration of WASI to allow plugins to access underlying operating system features such as the file system or networking.

6 Results & discussion

A technology comparison of 3 plugin systems of popular editors and IDEs and a plugin system technology was conducted. It revealed significant differences between the five criteria performance, plugin size, plugin isolation, plugin portability and plugin language interoperability, with no single technology standing out drastically across all dimensions. WebAssembly (Wasm) was then evaluated against the same criteria and put up for comparison in a technology comparison matrix shown in [Table 2](#). It performed at least equally well as the others in three out of the five criteria, only falling behind with slightly worse performance and plugin size compared to native machine code-based plugin systems. Despite these shortcomings, Wasm can offer significantly better plugin isolation through sandboxing and slightly better plugin portability. When compared to the plugin systems of the popular Visual Studio Code and IntelliJ-based IDEs, Wasm shows potential to outperform them for all criteria. The viability of Wasm for a text editor plugin system was proven in a proof of concept, where a Wasm plugin system was developed for the Helix editor. It required basic plugin functionalities such as the interaction between Wasm plugins and the core editor. These requirements were briefly reviewed, the criteria for good plugin systems were re-evaluated and other useful features such as a common interface definition or automatic binding generation demonstrated.

Wasm-based plugin systems present trade-offs: In comparison to native machine code-based plugin systems, they are slower and have larger plugins, however they can provide safety, portability and interoperability. Considering this, Wasm might be a better technology choice than traditional plugin systems, which rely on Java bytecode or JavaScript for plugins. Wasm beats the performance of these plugin systems, while also providing better plugin isolation and interoperability.

However challenges for Wasm as a plugin system technology remain. While Wasm itself is stable, its ecosystems consisting of proposals for new features and projects such as the Component Model have not reached stable versions yet. These technologies and projects are still rapidly evolving, leading to instabilities and frequent breaking changes between versions. Even though challenges for the Wasm ecosystem exist, multiple software projects, e.g. those presented in [Section 4.3](#), have already proven that Wasm is indeed ready for use in production-grade plugin systems.

7 Outlook

Looking back at the technology comparison, a set of five criteria was selected to evaluate the technologies under consideration. However, in retrospect, it became clear that these criteria were not sufficient for a meaningful evaluation. Instead more criteria regarding interface complexity, developer experience or the plugin system size should have also been considered. Also some criteria, such as the performance or plugin size have not been evaluated quantitatively, even though such quantitative analysis can lead to more accurate, objective and reproducible results. Furthermore, four plugin system technologies were analyzed, which is also not sufficient and representative of the current state of the art of plugin systems. Thus a larger variety of different plugin systems could have been evaluated.

Looking ahead, future directions include continuing the development of the proof of concept plugin system into a functional and usable plugin system for the Helix editor. The Wasm ecosystem itself is evolving, with exciting developments such as the Component Model or new proposals adding new features to Wasm. These developments in Wasm's ecosystem will continue to gain attention from researches and developers in various non-web areas such as avionics, automotive or distributed computing. The development of WebAssembly, driven by major web browsers, signifies its enduring presence and relevance both in web and possibly also non-web areas.

Bibliography

- [1] A. Haas *et al.*, “Bringing the Web up to Speed with WebAssembly,” in *ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI) 2017*, Association for Computing Machinery, Jun. 2017, pp. 185–200. doi: 10.1145/3062341.3062363.
- [2] “Qt Group Documentation.” Accessed: Apr. 21, 2025. [Online]. Available: <https://doc.qt.io/>
- [3] S. Padmanabhan and P. Jha, “WebAssembly at eBay: A Real-World Use Case.” Accessed: Apr. 21, 2025. [Online]. Available: <https://innovation.ebayinc.com/tech/engineering/webassembly-at-ebay-a-real-world-use-case/>
- [4] A. S. Tanenbaum and T. Austin, *Structured Computer Organization*, 6th ed. Pearson Education, Inc., 2025.
- [5] W. Zaeske, S. Friedrich, T. Schubert, and U. Durak, “WebAssembly in Avionics: Decoupling Software from Hardware,” in *IEEE/AIAA 42nd Digital Avionics Systems Conference (DASC)*, Nov. 2023. doi: 10.1109/dasc58513.2023.10311207.
- [6] M. N. Hoque and K. A. Harras, “WebAssembly for Edge Computing: Potential and Challenges,” in *IEEE Communications Standards Magazine*, IEEE, 2022, pp. 68–73. doi: 10.1109/MCOMSTD.0001.2000068.
- [7] S. Wallentowitz, B. Kersting, and D. M. Dumitriu, “Potential of WebAssembly for Embedded Systems,” in *11th Mediterranean Conference on Embedded Computing (MECO)*, IEEE, Jun. 2022. doi: 10.1109/meco55406.2022.9797106.
- [8] WebAssembly Community Group and A. Rossberg, “WebAssembly Core Specification (Release 2.0, Draft 2025-01-28).” Accessed: Mar. 27, 2025. [Online]. Available: <https://webassembly.github.io/spec/core>
- [9] A. Jangda, B. Powers, A. Guha, and E. Berger, “Mind the Gap: Analyzing the Performance of WebAssembly vs. Native Code,” Jan. 2019. doi: arXiv.1901.09056.
- [10] J. Bosamiya, W. S. Lim, and B. Parno, “Provably-Safe Multilingual Software Sandboxing using WebAssembly,” in *31st USENIX Security Symposium (USENIX Security 22)*, USENIX Association, Aug. 2022. Accessed: Apr. 15, 2025. [Online]. Available: <https://www.usenix.org/conference/usenixsecurity22/presentation/bosamiya>

-
- [11] B. Spies and M. Mock, “An Evaluation of WebAssembly in Non-Web Environments,” in *2021 XLVII Latin American Computing Conference (CLEI)*, 2021. doi: 10.1109/CLEI53233.2021.9640153.
- [12] Bytecode Alliance, “The WebAssembly Component Model.” Accessed: Apr. 16, 2025. [Online]. Available: <https://component-model.bytecodealliance.org/>
- [13] Bytecode Alliance, “component-model.” Accessed: Apr. 16, 2025. [Online]. Available: <https://github.com/WebAssembly/component-model>
- [14] “WASI.” Accessed: Apr. 17, 2025. [Online]. Available: <https://github.com/WebAssembly/WASI>
- [15] M. M. Syeed, A. Lokhman, T. Mikkonen, and I. Hammaouda, “Pluggable Systems as Architectural Pattern: An Ecosystemability Performance,” in *Proceedings of the 2015 European Conference on Software Architecture Workshops*, Association for Computing Machinery, 2015. doi: 10.1145/2797433.2797477.
- [16] “Linux man pages online.” Accessed: Apr. 16, 2025. [Online]. Available: <https://www.man7.org/linux/man-pages>
- [17] M. van Steen and A. S. Tanenbaum, *Distributed Systems: Principles and Paradigms*, Edition 4, version 03. Maarten van Steen, 2025. Accessed: Mar. 27, 2025. [Online]. Available: <https://www.distributed-systems.net/index.php/books/ds4/>
- [18] “StackOverflow Developer Survey 2024.” Accessed: Apr. 08, 2025. [Online]. Available: <https://survey.stackoverflow.co/2024/>
- [19] “Notepad++ User Manual.” Accessed: Apr. 16, 2025. [Online]. Available: <https://npp-user-manual.org/>
- [20] “VST 3 Developer Portal.” Accessed: Apr. 09, 2025. [Online]. Available: https://steinbergmedia.github.io/vst3_dev_portal
- [21] “Microsoft Flight Simulator - SDK documentation.” Accessed: Apr. 16, 2025. [Online]. Available: <https://docs.flightsimulator.com/>
- [22] “Electron Docs.” Accessed: Apr. 16, 2025. [Online]. Available: <https://www.electronjs.org/docs/latest>
- [23] “Visual Studio Code documentation.” Accessed: Apr. 13, 2025. [Online]. Available: <https://code.visualstudio.com/docs>
- [24] “Zed documentation.” Accessed: Apr. 19, 2025. [Online]. Available: <https://zed.dev/docs/>

-
- [25] “ZelliJ User Guide.” Accessed: Apr. 19, 2025. [Online]. Available: <https://zellij.dev/documentation/>
- [26] “Extism documentation.” Accessed: Apr. 19, 2025. [Online]. Available: <https://extism.org/docs/overview>